

DECLARAÇÃO DO(A) ORIENTADOR(A)
(Concordância com a Composição da Banca Examinadora)

UNIVERSIDADE FEDERAL DO CEARÁ
Pró-Reitoria de Pesquisa e Pós-Graduação
Programa de Pós-Graduação em Ciências da
Computação

DECLARAÇÃO DE CONCORDÂNCIA DO(A) ORIENTADOR(A) COM A COMPOSIÇÃO DA BANCA EXAMINADORA

Eu, César Lincoln Cavalcante Mattos, professor(a) vinculado(a) ao Programa de Pós-Graduação em Ciências da Computação da Universidade Federal do Ceará, matrícula SIAPE nº 1407874, declaro, para os devidos fins, que concordo com a composição da banca examinadora sugerida para a defesa de Proposta de mestrado do(a) discente Hugo Ramos Soares, matrícula nº 564305, a ser realizada no dia dez de julho de 2026.

A banca será composta pelos seguintes membros:

Presidente (Orientador(a)):

Prof. Dr. César Lincoln Cavalcante Mattos – Universidade Federal do Ceará (UFC)

Membros Titulares:

Prof. Dr. João Paulo Pordeus Gomes – Universidade Federal do Ceará (UFC)

Prof. Dr. José Antônio Fernandes de Macêdo – Universidade Federal do Ceará (UFC)

Prof. Dr. Diego Parente Paiva Mesquita – Fundação Getúlio Vargas (FGV)

Declaro, ainda, que a composição acima atende aos critérios estabelecidos pelo Programa de Pós-Graduação em Ciências da Computação, bem como às normas da Universidade Federal do Ceará e da CAPES para a realização de defesas acadêmicas no âmbito da pós-graduação stricto sensu.

Por ser verdade, firmo a presente declaração.

Fortaleza, 23 de Junho de 2026.

Prof. Dr. César Lincoln Cavalcante Mattos
Orientador(a)
Programa de Pós-Graduação em Ciências da Computação
Universidade Federal do Ceará

DECLARAÇÃO SOBRE O USO OU NÃO USO DE INTELIGÊNCIA ARTIFICIAL

Eu, Hugo Ramos Soares, responsável pela obra HoTHP: Hyperbolic Rotary Position Embedding-based Transformer Hawkes Process for Length Extrapolation, declaramos que as informações prestadas refletem de forma verdadeira e completa o uso ou não uso de ferramentas de Inteligência Artificial nesta produção, em conformidade com a Resolução de Integridade da UFC e a PORTARIA Nº 39/PRPPG/UFC, DE 01 DE OUTUBRO DE 2025.

SELECIONE UMA DAS OPÇÕES:

☐ NÃO HOUVE USO DE FERRAMENTAS DE INTELIGÊNCIA ARTIFICIAL

Declaro que não utilizei, em nenhuma etapa do desenvolvimento do presente trabalho, ferramentas de Inteligência Artificial, incluindo mas não se limitando a recursos de geração de texto, imagem, código, resumos, traduções ou análises. Assumo integral responsabilidade pelo conteúdo, conforme os princípios de integridade acadêmica e a legislação vigente.

(Se marcar esta opção, não é necessário preencher os campos 1 a 5 abaixo.)

☒ HOUVE USO DE FERRAMENTAS DE INTELIGÊNCIA ARTIFICIAL

(Preencher os campos obrigatórios abaixo.)

1. Ferramenta(s) e versão(ões) predominante(s): Claude Opus 4.6, 4.7 e 4.8

2. Período de uso: Junho/2025 a Junho/2026

3. Finalidade(s) (marcar todas que se aplicam):

☒ Exploração inicial de ideias

☐ Busca/triagem de literatura

☐ Leitura assistida/Resumos (com conferência humana)

☐ Revisão linguística/Referências

☒ Apoio reprodutível à análise/apresentação (restrito à geração automatizada e reprodutível de tabelas, gráficos, figuras ou visualizações a partir de dados previamente analisados e interpretados pelos autores, sem substituição da análise científica humana)

☒ Programação (sugestão/depuração/documentação)

☐ Transcrição (com anonimização): conversão literal, mediante autorização, de áudios ou vídeos em texto de entrevistas, aulas, reuniões ou palestras, com obrigatória revisão humana, anonimização e respeito aos direitos autorais, vedada a identificação de voz ou outros dados biométricos.

☒ Tradução técnica (com revisão humana)

4. Descrição sintética do uso / prompts-tipo: Usei o Claude para discutir conceitos de arquiteturas de Transformer Hawkes Processes (THP, RoTHP, kernel hiperbólico vs. trigonométrico), gerar scripts de tabelas, gráficos e figuras a partir de resultados que eu já tinha obtido e interpretado, sugerir, debugar e documentar código e traduzir o abstract e trechos técnicos para o inglês. Prompts como "explique a diferença entre kernel trigonométrico e hiperbólico", "gere um gráfico comparando essas métricas", "traduza este parágrafo técnico".

5. Validação humana (checagens, testes, leitura crítica): Conferi as explicações conceituais com a literatura, executei e testei todo o código sugerido, reproduzindo os resultados, verifiquei figuras e

tabelas e revisei as traduções linha a linha. A análise, interpretação e conclusões são minhas e assumo responsabilidade integral pelo conteúdo.

Declarações éticas (obrigatórias caso tenha havido uso de IA):

Todas as opções abaixo devem ser obrigatoriamente assinaladas para que o trabalho possa ser submetido à defesa/apresentação. O não preenchimento completo impedirá a continuidade do processo.

- ☒ Não houve geração de conteúdo original, ideias, interpretações ou análises pela IA;
- ☒ Não enviei dados inéditos ou sensíveis à serviços que utilizam conteúdos para treinamento de modelos, exceto em plataformas institucionais ou com garantias contratuais de confidencialidade e não retenção, assegurando conformidade com a LGPD (Lei nº 13.709/2018) e demais normas de proteção de dados;
- ☒ Respeitei direitos autorais, licenças, confidencialidade e políticas editoriais;
- ☒ Em transcrições, apliquei anonimização e não realizei identificação por voz/biometria;
- ☒ Assumo responsabilidade integral e exclusiva pelo conteúdo final desta obra.

Data: 23 de junho de 2026

Nome do Autor: Hugo Ramos Soares

Assinatura:  Documento assinado digitalmente
HUGO RAMOS SOARES
Data: 24/06/2026 08:59:35-0300
Verifique em <https://validar.iti.gov.br>

NOME DO DOCUMENTO

**Proposta de Dissertação - Hugo Ramos.
pdf**

NÚMERO DE PALAVRAS

18146 Words

NÚMERO DE CARACTERES

89897 Characters

NÚMERO DE PÁGINAS

59 Pages

TAMANHO DO ARQUIVO

2.5MB

DATA DE ENVIO

Jun 23, 2026 8:11 PM GMT-3

DATA DO RELATÓRIO

Jun 23, 2026 8:12 PM GMT-3**● 16% geral de similaridade**

O total combinado de todas as correspondências, incluindo fontes sobrepostas, para cada banco de dados.

- 13% Banco de dados da Internet
- Banco de dados do Crossref
- 11% Banco de dados de trabalhos enviados
- 11% Banco de dados de publicações
- Banco de dados de conteúdo publicado no Crossref



16

UNIVERSIDADE FEDERAL DO CEARÁ
CENTRO DE CIÊNCIAS
DEPARTAMENTO DE COMPUTAÇÃO
PROGRAMA DE PÓS-GRADUAÇÃO EM CIÊNCIA DA COMPUTAÇÃO
MESTRADO ACADÊMICO EM LÓGICA E INTELIGÊNCIA ARTIFICIAL

HUGO RAMOS SOARES

1

HOTHP: HYPERBOLIC ROTARY POSITION EMBEDDING-BASED TRANSFORMER
HAWKES PROCESS FOR LENGTH EXTRAPOLATION

FORTALEZA

2026

HUGO RAMOS SOARES

¹ HOTHF: HYPERBOLIC ROTARY POSITION EMBEDDING-BASED TRANSFORMER
HAWKES PROCESS FOR LENGTH EXTRAPOLATION

³⁴ Dissertação apresentada ao Curso de Mestrado Acadêmico em Lógica e Inteligência Artificial do Programa de Pós-Graduação em Ciência da Computação do Centro de Ciências da Universidade Federal do Ceará, como requisito parcial à obtenção do título de mestre em Ciência da Computação. Área de Concentração: Ciência da Computação

⁷ Orientador: Prof. Dr. César Lincoln Cavalcante Mattos

Coorientador: Prof. Dr. João Paulo Pordeus Gomes

FORTALEZA

2026

HUGO RAMOS SOARES

HOTHP: HYPERBOLIC ROTARY POSITION EMBEDDING-BASED TRANSFORMER
HAWKES PROCESS FOR LENGTH EXTRAPOLATION

34
7
Dissertação apresentada ao Curso de Mestrado Acadêmico em Lógica e Inteligência Artificial do Programa de Pós-Graduação em Ciência da Computação do Centro de Ciências da Universidade Federal do Ceará, como requisito parcial à obtenção do título de mestre em Ciência da Computação. Área de Concentração: Ciência da Computação

Aprovada em:

BANCA EXAMINADORA

Prof. Dr. César Lincoln Cavalcante
Mattos (Orientador)
Universidade Federal do Ceará (UFC)

Prof. Dr. João Paulo Pordeus Gomes (Coorientador)

16
Prof. Dr. José Antônio Fernandes de Macêdo
Universidade Federal do Ceará (UFC)

Prof. Dr. Diego Parente Paiva Mesquita
Fundação Getúlio Vargas (FGV)

“Não é porque as coisas são difíceis que não ousamos; é porque não ousamos que elas são difíceis.”

(Sêneca)

RESUMO

Processos Pontuais Temporais (*Temporal Point Process* (TPP)) baseados em Transformers têm apresentado resultados satisfatórios na modelagem de sequências de eventos. No entanto, esses modelos frequentemente degradam quando testados em sequências mais longas do que aquelas enfrentadas durante o treinamento, o que se trata de uma importante limitação ao abordarmos aplicações reais.

Neste trabalho, propomos o HoTHP, um modelo baseado em Transformer para Processos de Hawkes que substitui o kernel rotativo trigonométrico do RoTHP por uma versão utilizando funções hiperbólicas, inspirado no HoPE. Ao invés de rotacionar os vetores Q e K antes do produto da atenção, o HoTHP reformula o score de atenção como uma função da diferença temporal entre os eventos, combinando o kernel hiperbólico com um parâmetro de decaimento θ' , aprendido durante o treinamento e restrito a valores maiores que o maior valor de θ , para garantir o decaimento monotônico.

Esse modelo incorpora um viés de decaimento explícito no mecanismo de atenção, fazendo com que eventos recentes recebam maior peso na construção do estado oculto. Diferentemente do que se poderia supor, esse viés de recência atua exclusivamente na atenção e não na função de intensidade, permitindo que o HoTHP modele tanto processos auto-excitatórios quanto auto-inibitórios.

Avaliamos o HoTHP em dados sintéticos gerados por processos de Hawkes e em quatro datasets reais (Amazon, Stackoverflow, Taxi e Retweet), seguindo o protocolo “train-short, test-long”, no qual o modelo é treinado em sequências curtas e testado em sequências até $5\times$ mais longas. Nos dados sintéticos, o HoTHP mantém estável a negativa do log da verossimilhança (*Negative Log-Likelihood* (NLL)) mesmo em fatores de extrapolação elevados, enquanto o RoTHP degrada progressivamente. Nos dados reais, o HoTHP obtém a melhor NLL em extrapolação no Amazon e Taxi, empata no Stackoverflow, e supera o RoTHP no Retweet após normalização temporal, com NLL $5\times$ de $-0,542 \pm 0,004$ contra $88,9 \pm 144,8$ do RoTHP. O custo computacional adicional do HoTHP no treinamento ($1,05\times$ a $2,24\times$) é pago uma única vez, enquanto a estabilidade na extrapolação beneficia toda inferência subsequente.

Palavras-chave: Processos Pontuais Temporais. Processo de Hawkes. Transformers. Extrapolação de comprimento.

ABSTRACT

Transformer-based Temporal Point Processes (TPP) have shown promising results in event sequence modeling. However, these models frequently degrade¹ when tested on sequences longer than those encountered during training, which represents an important limitation when addressing real-world applications.

⁷⁵ In this work, we propose HoTHP, a Transformer-based model for Hawkes Processes that replaces the trigonometric rotary kernel of RoTHP with a version using hyperbolic functions, inspired by HoPE. Instead of rotating the Q and K vectors before the attention product, HoTHP reformulates the attention score as a function of the temporal difference between events, combining the hyperbolic kernel with a learned decay parameter θ' , constrained during training to values greater than the largest θ value, in order to guarantee monotonic decay.

This model incorporates an explicit decay bias into the attention mechanism, assigning greater weight to recent events when building the hidden state. Contrary to what one might expect, this recency bias acts exclusively on the attention mechanism and not on the intensity function, allowing HoTHP to model both self-exciting and self-inhibiting processes.

We evaluate HoTHP on synthetic datasets generated by Hawkes processes and⁴⁵ on four real-world datasets (Amazon, Stackoverflow, Taxi, and Retweet), following the “train-short, test-long” protocol, in which the model is trained on short sequences and tested on sequences up to $5\times$ longer. On synthetic data, HoTHP maintains a stable negative log-likelihood (NLL) even at high extrapolation factors, while RoTHP degrades progressively. On real data, HoTHP achieves the best extrapolation NLL on Amazon and Taxi, ties on Stackoverflow, and outperforms RoTHP on Retweet after temporal normalization, with $5\times$ NLL of -0.542 ± 0.004 versus 88.9 ± 144.8 for RoTHP. The additional training cost of HoTHP ($1.05\times$ to $2.24\times$) is paid once, while the extrapolation stability benefits all subsequent inference.

Keywords: Temporal Point Process. Hawkes Process. Transformers. Length extrapolation.

Figura 1	– Poisson process.	16
Figura 2	– Probability intensity function	19
Figura 3	– Attention x Multi-Head Attention.	24
Figura 4	– Architecture of the Transformer Hawkes Process	26
Figura 5	– Rotary Position Embeddings (RoPE).	28
Figura 6	– As the extrapolation increases, we clearly see that the degradation shown by RoPE is smaller than that of traditional sinusoidal methods, but still much sharper than the model proposed by Press <i>et al.</i> (2021)	30
Figura 7	– teste hugo	34
Figura 8	– Effect of the exponential decay on the hyperbolic functions ($\theta = 1,0$, $\theta' = 1,5$). Left: separate components, showing that cosh and sinh grow but the decay falls. Center: the product results in functions that decay monotonically. Right: decomposition into pure exponentials, both with a negative exponent, ensuring numerical stability.	36
Figura 9	– Comparison of the attention kernels as a function of the temporal distance Δt . In the training zone, both models can fit the data. In the extrapolation zone, the trigonometric kernel of the RoTHP oscillates in phases that were not learned, while the hyperbolic kernel of the HoTHP decays monotonically, like the classical Hawkes kernel.	40
Figura 10	– NLL (mean $\pm$ standard deviation, 10 seeds) as a function of the extrapolation factor α in the synthetic scenarios. The HoTHP keeps the NLL more stable as α grows, while the THP degrades drastically and the RoTHP and NHP show intermediate degradation.	45
Figura 11	– RMSE of the next event prediction (mean $\pm$ standard deviation, 10 seeds) as a function of the extrapolation factor α . Unlike the NLL, the RMSE does not separate the models clearly: in the slow decay all of them degrade in a similar way and in the fast decay they stay close to each other.	46

Figura 12 – NLL (mean $\pm$ standard deviation, 3 seeds) as a function of the extrapolation factor on the real datasets, for the three Transformer models. The NHP is omitted from the plot for readability, since its lower NLL on Taxi and Amazon would compress the scale; its values are in Table 6. Among the Transformer models, the HoTHP keeps the lowest and most stable NLL in extrapolation.	50
Figura 13 – Type-prediction accuracy (mean $\pm$ standard deviation, 3 seeds) as a function of the extrapolation factor on the real datasets, for the three Transformer models (the NHP is in Table 7). On Amazon, the HoTHP shows the clearest and growing advantage as the sequence gets longer.	50
Figura 14 – Convergence curves of the validation NLL (mean $\pm$ standard deviation, 3 seeds). The first epochs were omitted to improve the visualization. On Amazon and Stackoverflow, all the models converge smoothly. On Taxi, the RoTHP is more unstable in the early epochs. On Retweet, the wide band of the RoTHP reflects the collapse of 2 of the 3 seeds.	51
Figura 15 – Conditional intensity $\lambda^*(t)$ learned by the THP (orange), RoTHP (blue), and HoTHP (green) for one test sequence of each dataset. The gray vertical bars indicate the events. On Amazon, the intensity grows between events and falls when an event occurs (self-inhibition). On Stackoverflow, the intensity is approximately constant (memoryless process). On Retweet, both models collapse to $\lambda^* \approx 0$ after the initial burst due to the extreme temporal scale. .	52
Figura 16 – NLL on the normalized Retweet (mean $\pm$ standard deviation, 3 seeds). On the left, the $1 \times$ condition (training length): both models obtain a similar NLL. On the right, the $5 \times$ condition (extrapolation): the HoTHP keeps a stable NLL, while the RoTHP shows high variance due to one seed that diverges. .	54

LISTA DE TABELAS

Tabela 1	– Model log-likelihood by dataset	29
Tabela 2	– Accuracy Comparison	29
Tabela 3	– RMSE Comparison	29
Tabela 4	– NLL (mean $\pm$ standard deviation, 10 seeds) on the synthetic data. Lower values indicate better performance. The best result for each configuration is in bold.	44
Tabela 5	– Real datasets used in the experiments. L_{train} is the maximum training length (number of events) and $L_{5\times}$ is the test length in the extrapolation condition. .	47
Tabela 6	– NLL on the real datasets (mean $\pm$ standard deviation, 3 seeds), for the extrapolation factors $1\times$, $2\times$, and $5\times$. Lower is better. The best value in each column is in bold. NHP is a recurrent baseline (CT-LSTM); the other three share the same Transformer architecture and intensity function and differ only in the positional kernel.	48
Tabela 7	– Type-prediction accuracy on the real datasets (mean $\pm$ standard deviation, 3 seeds), for the factors $1\times$, $2\times$, and $5\times$. Higher is better. The best value in each column is in bold.	48
Tabela 8	– Ablation: effect of the temporal normalization on Retweet. “Original (raw)” uses the raw timestamps; “Normalized” uses the timestamps divided by the mean of the gaps. Mean $\pm$ standard deviation (3 seeds). The best value in each column of the normalized block is in bold.	53
Tabela 9	– Average training time in minutes (mean over 3 seeds, except NHP on Amazon with 2 seeds; NHP was not run on the normalized Retweet). All the models use the same architecture size and training budget (100 epochs).	55

LISTA DE ABREVIATURAS E SIGLAS

¹ HoPE	<i>Hyperbolic Rotary Positional Encoding for Stable Long-Range Dependency Modeling in Large Language Models</i>
HoTHP	<i>Hyperbolic Transformer Hawkes Process</i> ¹⁰⁴
LSTM	<i>Long Short-Term Memory</i>
NHP	<i>Neural Hawkes Process</i>
⁹¹ NLL	<i>Negative Log-Likelihood</i>
RMSE	<i>Root Mean Squared Error</i>
RoTHP	<i>Rotary Transformer Hawkes Process</i> ¹⁰⁷
SAHP	<i>Self-Attentive Hawkes Process</i>
THP	<i>Transformer Hawkes Process</i>
TPP	<i>Temporal Point Process</i>

SUMÁRIO

1	INTRODUCTION	12
1.1	Objectives	13
1.2	Contributions	13
1.3	Text organization	13
2	THEORETICAL BACKGROUND	14
2.1	Poisson distribution	14
2.2	Poisson processes	15
73 2.3	Temporal point processes	16
2.3.1	<i>Loss function</i>	17
2.4	Hawkes processes	18
2.4.1	<i>Relationship with the Cox process</i>	20
2.5	Neural temporal point processes	20
2.6	Transformers and the attention mechanism	22
2.6.1	<i>Encoders and Decoders</i>	23
71 2.6.2	Attention mechanism	23
2.6.3	Multi-Head Attention	24
3	RELATED WORK	25
3.1	Transformer Hawkes Process (THP)	25
3.2	Rotary Position Embeddings (RoPE) and the RoTHP	27
3.3	Length Extrapolation and inductive bias	30
3.3.1	<i>Attention with Linear Biases (ALiBi)</i>	31
3.3.2	<i>Inductive bias and its relevance to the Hawkes process</i>	31
4	METHODOLOGY	33
4.1	The oscillatory limitation of the RoTHP	33
4.2	From the trigonometric kernel to the hyperbolic kernel	34
4.3	Exponential decay and the parameter θ'	35
4.3.1	<i>Parametrization of the learnable θ'</i>	35
4.4	Attention score of the HoTHP	36
4.5	Numerical stability	37
4.6	Temporal normalization	37

4.7	Complete architecture	38
4.8	Where the exponential bias acts	39
4.9	Experimental setup	41
4.9.1	<i>Synthetic data generation</i>	41
4.9.2	<i>Train short, test long protocol</i>	41
4.9.3	<i>Compared models and implementation</i>	42
4.9.4	<i>Metrics and statistical evaluation</i>	42
4.9.5	<i>Computational environment</i>	43
5	RESULTS	44
5.1	Synthetic data	44
5.2	Real data	46
5.2.1	<i>Training convergence</i>	50
5.3	Intensities learned by the models	51
5.4	Ablation: temporal normalization of Retweet	53
5.5	Discussion	54
5.5.1	<i>Computational cost versus benefit in extrapolation</i>	55
	REFERÊNCIAS	57

1 INTRODUCTION

Many real-world phenomena can be modeled as sequences of events that happen at irregular instants over time, such as financial transactions, earthquakes, interactions in social networks, neuron firings, and so on. ¹ Temporal point processes (TPPs) are the mathematical framework for modeling these sequences. The central goal is to estimate the conditional intensity function $\lambda^*(t)$, which describes the instantaneous rate of event occurrence given the history. ⁴³

Among the classical models, the Hawkes process (HAWKES, 1971) stands out for capturing self-excitation, that is, when ⁴ the occurrence of an event temporarily increases the probability of future events, with the influence decaying exponentially over time. More recent neural models, such as ⁴ the Transformer Hawkes Process (*Transformer Hawkes Process* (THP)) (ZUO *et al.*, 2021), replace the parametric kernel with an attention mechanism that learns the relationship between events directly from the data. The *Rotary Transformer Hawkes Process* (RoTHP) (DAI *et al.*, 2026) moves further in this direction by using ⁷⁰ rotary positional embeddings (RoPE) (SU *et al.*, 2023) to encode relative temporal distances in the attention mechanism.

An open problem in these models is length extrapolation. ¹⁰⁸ When they are trained on short sequences and tested on longer ones, performance tends to degrade. In the THP, this happens because the absolute temporal encoding produces representations that fall outside the training distribution for new positions. In the RoTHP, the trigonometric kernel (based on sine and cosine) is oscillatory, and for temporal distances not seen during training the attention score can take arbitrary values, with no guarantee of decay.

This work proposes the *Hyperbolic Transformer Hawkes Process* (HoTHP), which replaces the trigonometric kernel of the RoTHP with a hyperbolic kernel that has exponential decay, inspired by the ¹ *Hyperbolic Rotary Positional Encoding for Stable Long-Range Dependency Modeling in Large Language Models* (HoPE) (DAI *et al.*, 2025). The idea is to embed a recency bias into the attention mechanism, so that recent events receive larger weights and distant events receive weights that tend monotonically to zero. This bias is structural (it does not depend on the training data) and holds for arbitrarily large temporal distances, which makes extrapolation possible.

57

1.1 Objectives

The general objective of this work is to investigate whether replacing the oscillatory kernel with a kernel that has exponential decay in the attention mechanism improves the length extrapolation ability of neural models for temporal point processes. To do this, we propose and implement the HoTHP, adapting the HoPE (DAI *et al.*, 2025) from discrete positions to continuous instants. We evaluate the model with the *train short, test long* technique on synthetic data, where the generating process is known, and on real data from varied domains. Finally, we analyze under which conditions the recency bias is beneficial and under which it is neutral or harmful, through the visualization of the learned intensities, accuracy, ¹Root Mean Squared Error (RMSE), and ablations on the temporal scale of the data.

1.2 Contributions

This work contributes the proposal of the HoTHP, which introduces an exponential recency bias directly into the temporal attention kernel, with a decay parameter θ' that is learnable and structurally constrained to guarantee monotonic decay. The experimental analysis reveals that this bias acts on the construction of the hidden state (through attention) and not on the intensity function, which allows the model to learn both self-exciting and self-inhibiting processes, as long as recent events are more informative than distant ones. The evaluation on synthetic data and four real datasets identifies the conditions that determine when the bias is beneficial and when it is not (processes without memory or with an extreme temporal scale).

44

1.3 Text organization

The remainder of this dissertation is organized as follows. Chapter 2 presents the theoretical background on point processes, Hawkes processes, neural Hawkes processes (NHP), and the Transformer architecture with its attention mechanism. Chapter 3 reviews the related work: the Transformer-based approaches for point processes (THP and RoTHP) and the length extrapolation and inductive bias techniques (RoPE, ALiBi) that underlie the HoTHP. Chapter 4 describes the HoTHP architecture, the formulation of the hyperbolic kernel, and the experimental setup. Chapter 5 presents ⁹⁰the results on synthetic and real data, the learned intensities, and the temporal normalization ablation. Chapter ?? concludes the work and proposes future directions.

2 THEORETICAL BACKGROUND

Understanding the behavior of event sequences over time and how events influence one another requires the use of complex mathematical tools. This chapter presents the theoretical background needed to support the proposed model. We start with the Poisson distribution and the modeling of point processes, which are needed to understand how events happen in continuous time in an irregular way. Next, we look at the Hawkes process, which shows how events can affect the probability of future events. After that, we present neural Hawkes processes, which replace the parametric kernel with neural networks, culminating in the *Neural Hawkes Process* (NHP), and we review the transformer architecture and its attention mechanism. The Transformer-based Hawkes process approaches (THP and RoTHP) and the length extrapolation techniques, which are the direct basis of the proposed model, are covered in Chapter 3.

2.1 Poisson distribution

Before dealing with the Poisson process, it is worth recalling the Poisson distribution, from which it takes its name. The Poisson distribution is a discrete probability distribution that models the number of events that occur in a fixed interval of time or space, under the assumptions that the events are independent and occur at a constant average rate λ . The probability of observing exactly k events is given by:

$$\Pr(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}, \quad k = 0, 1, 2, \dots, \quad (2.1)$$

where $\lambda > 0$ is the average number of events expected in the interval. Both the mean and the variance of this distribution are equal to λ , so that a higher rate implies both more expected events and greater variability around that expectation.

The Poisson distribution can also be obtained as an approximation to the binomial distribution. If X counts the number of successes in n independent trials, each with success probability p , and we write $\lambda = np$, then when n is large and p is small enough that np is of moderate size, the binomial probabilities are well approximated by the expression in Equation 2.1 (ROSS, 2019). This justifies the use of the Poisson distribution to count independent occurrences when each one is individually unlikely but a large number of opportunities is available.

The same distribution underlies the Poisson process presented in the next section, which generates events over time so that the count in any interval is Poisson distributed. There,

Equation 2.4 reappears with $\Lambda = \lambda T$ in place of λ , so that the expected number of events grows proportionally to the length T of the observed interval.

2.2 Poisson processes

The Poisson process is a mathematical object that consists of random points in a mathematical space with the important feature of no dependence between the points. This model is widely used in areas such as queueing theory, astronomy, biology, and survival analysis. The name comes from the fact that the points, or events, occur following a Poisson distribution.

The Poisson process is the most common form of point process. Daley e Vere-Jones (2003) define a stationary Poisson process by the following equation, where we use $N(a_i, b_i]$ to represent the number of events of the process.

$$\Pr\{N((a_i, b_i]) = n_i, i = 1, \dots, k\} = \prod_{i=1}^k \frac{[\lambda(b_i - a_i)]^{n_i}}{n_i!} e^{-\lambda(b_i - a_i)}. \quad (2.2)$$

This definition covers three important aspects: (i) the number of points in each finite interval $(a_i, b_i]$ has a Poisson distribution; (ii) the numbers of points in disjoint intervals are independent random variables; and (iii) the distributions are stationary, that is, they depend only on the lengths $b_i - a_i$ of the intervals. The constant λ is the average rate or the average density of the events of the process. In the following subsections we will see that λ is also the intensity with which the events occur.

Considering the special case of a single interval $k = 1$, we have $(a_i, b_i] = (0, T]$. Then, the interval size is $b_i - a_i = T$. Substituting into Equation 2.2, we have:

$$\Pr\{N((0, T]) = n_1\} = \frac{[\lambda T]^{n_1}}{n_1!} e^{-\lambda T}. \quad (2.3)$$


In Equation 2.3, replacing $(0, T]$ by N and λT by Λ , we obtain the classical form of the Poisson distribution associated with the Poisson process:

$$\Pr(N = n) = \frac{\Lambda^n}{n!} e^{-\Lambda}. \quad (2.4)$$

Equation 2.4 makes explicit the relationship between the Poisson distribution and the Poisson process: the number of events observed in any interval of size T is Poisson with mean $\Lambda = \lambda T$. In other words, the Poisson process is a way of generating event times whose count satisfies the Poisson distribution and that are independent and stationary. Poisson processes are the most fundamental form of temporal point process, which we will discuss in the next section.

Figure 1 illustrates a homogeneous Poisson process in the interval $[0, T]$. The horizontal axis represents time and each vertical mark represents the occurrence of an event generated by the process. Although the events look irregular, they were generated so that the waiting time between one event and another is exponential $1/\lambda$ and independent, where λ is constant and therefore the expected number of events in a given interval T is defined by λT . This illustrates the intuition behind the Poisson process, which is the random occurrence of events, independent and following a constant rate λ .

Figura 1 – Poisson process.



2.3 Temporal point processes

Temporal point processes (TPP) are probabilistic models designed to describe sequences of discrete events that happen in continuous time (ZUO *et al.*, 2021). Unlike discrete models, which assume that events happen at known time intervals, TPPs make it possible to handle events that happen at irregular, arbitrary times. This makes them particularly useful in domains where the moment of the event occurrence is as important as the event itself.

Many scenarios generate such event sequences, from neuron spikes in neuroscience, to buying and selling stocks in the financial market, to tweets in social networks, etc. (ZHOU *et al.*, 2025). These event sequences, made up of asynchronous events, usually influence one another, making the study of these phenomena not only complex but also necessary.

TPPs are formally defined by a conditional intensity function $\lambda(t | \mathcal{H}_t)$, where $\mathcal{H}_t$ is the history of events up to time t . The intensity is the instantaneous rate at which an event is expected to occur at a time t , given the past events.

One variation of TPPs is the existence of *marks*, that is, event types, where the sequence is now described as $X = \{(t_1, m_1), \dots, (t_N, m_N)\}$, where N represents the number of

events and each m_i is the mark (type) associated with event i .

It is worth noting that the term *mark* is not used in the same way by every author. Part of the literature treats the mark as the event type, that is, a category from a finite set. This is the view of Du *et al.* (2016), who represent the mark as a one-hot vector and predict it as one class among K . Another part of the literature treats the mark in a broader sense, as any extra information attached to an event, which may be continuous. Rasmussen (2018), for example, describes the mark as the additional information associated with an event, such as the magnitude of an earthquake, and lets the mark space be a subset of $\mathbb{R}$ or of $\mathbb{N}$. In this work we adopt the first convention: the mark is always the event type, a categorical value, which is the discrete case ($\mathbb{N}$) of the broader definition. We therefore use the words *mark* and *type* interchangeably throughout the text.

The distribution of a TPP with K marks can be characterized by K conditional intensity functions $\lambda_k(t \mid \mathcal{H}_t)$ (one for each mark k) and is defined as

$$\lambda_k(t \mid \mathcal{H}_t) = \lim_{\Delta t \rightarrow 0} \frac{\Pr(\text{event of type } k \text{ in } [t, t + \Delta t) \mid \mathcal{H}_t)}{\Delta t}. \quad (2.5)$$

Note that the definition above can also be applied to TPPs without marks, simply by setting $K = 1$ (SHCHUR *et al.*, 2021).

2.3.1 Loss function

For point processes, given the observation of all events in the interval $[0, T]$, we use the log-likelihood as in the equation below:

$$\mathcal{L} = \sum_{i: t_i \leq T} \log \lambda_{k_i}(t_i) - \underbrace{\int_{t=0}^T \lambda(t) dt}_{\text{Let us call it } \Lambda}. \quad (2.6)$$

We will derive this function following Meier and Eisner (2017). First, we define $N(t) = |\{h : t_h \leq t\}|$ as the count of events (of any type) before time t . Given the past history $\mathcal{H}_i$, the number of events in $(t_{i-1}, t]$ is $\Delta N(t_{i-1}, t) \stackrel{\text{def}}{=} N(t) - N(t_{i-1})$.

Let $T_i > t_{i-1}$ be the random variable of the time of the next event and K_{i+1} the random variable of the type of the next event. The cumulative distribution function and the

probability density function of T_i (conditioned on $\mathcal{H}_i$) are given by:

$$F(t) = P(T_i \leq t) = 1 - P(T_i > t) \quad (2.7)$$

$$= 1 - P(\Delta N(t_{i-1}, t) = 0) \quad (2.8)$$

$$= 1 - \exp\left(-\int_{t_{i-1}}^t \lambda(s) ds\right) \quad (2.9)$$

$$= 1 - \exp(\Lambda(t_{i-1}) - \Lambda(t)) \quad (2.10)$$

$$f(t) = \exp(\Lambda(t_{i-1}) - \Lambda(t)) \lambda(t), \quad (2.11)$$

where $\Lambda(t) = \int_0^t \lambda(s) ds$ and $\lambda(t) = \sum_{k=1}^K \lambda_k(t)$.

In addition, given the past history $\mathcal{H}_i$ and the time of the next event t_i , the distribution of the event type k_i (probability of being of type k_i) is given by:

$$P(K_i = k_i | t_i) = \frac{\lambda_{k_i}(t_i)}{\lambda(t_i)} \quad (2.12)$$

Therefore, we can derive the likelihood function L as follows:

$$L = \prod_{i: t_i \leq T} L_i \quad (2.13)$$

$$= \prod_{i: t_i \leq T} \{f(t_i) P(K_i = k_i | t_i)\} \quad (2.14)$$

$$= \prod_{i: t_i \leq T} \{\exp(\Lambda(t_{i-1}) - \Lambda(t_i)) \lambda_{k_i}(t_i)\} \quad (2.15)$$

Finally, by applying the logarithm to the likelihood we turn the product into a sum:

$$\mathcal{L} \stackrel{\text{def}}{=} \log L \quad (2.16)$$

$$= \sum_{i: t_i \leq T} \log \lambda_{k_i}(t_i) - \sum_{i: t_i \leq T} (\Lambda(t_i) - \Lambda(t_{i-1})) \quad (2.17)$$

$$= \sum_{i: t_i \leq T} \log \lambda_{k_i}(t_i) - \Lambda(T) \quad (2.18)$$

$$= \sum_{i: t_i \leq T} \log \lambda_{k_i}(t_i) - \int_{t=0}^T \lambda(t) dt \quad (2.19)$$

2.4 Hawkes processes

Poisson processes are a basic way of working with sequences of discrete events; however, they start from the assumption that such events are independent of one another. Even in a multivariate and non-homogeneous Poisson process, where the variation in the probability of an event occurring at t varies as a function of t , the events are still independent of one another.

It assumes that an event of type k happens in the infinitesimal interval $[t, t + dt]$ with probability $\lambda_k(t)dt$. The total intensity of all events is given by $\lambda(t) = \sum_{k=1}^K \lambda_k(t)$.

The Hawkes process, on the other hand, starts from the principle that past events influence the probability of future events. In a traditional Hawkes process, we assume that this probability is always increased, added to the probability from past events, and decays exponentially (MEI; EISNER, 2017).


The Hawkes process is considered doubly stochastic (ZUO *et al.*, 2021), because $\lambda(t)$ is a random variable, as in a Cox process, but it is also a Poisson process. The intensity of a traditional univariate Hawkes process is given by:

$$\lambda(t | \mathcal{H}_t) = \mu + \sum_{t_i < t} \alpha e^{-\beta(t-t_i)}, \quad (2.20)$$

where α represents the instantaneous increase in the probability intensity caused by event e and β represents the speed at which this intensity decays over time.

Figure 2 illustrates the probability intensity λ of an event where, from $t = 0$ to $t = 2$, this intensity stays at a baseline μ . At time $t = 2$, event e happens, raising the probability intensity instantaneously. Then, λ decays exponentially, returning to the stationary state μ , which represents the absence of influence from any event at time t .

Figure 2 – Probability intensity function
Probability intensity function representing instantaneous self-excitation and exponential decay.



Throughout this work we will consider the multivariate Hawkes process, which is able to capture the influence of events of different types k on one another, such as a post on a social network positively influencing the occurrence of another type of interaction, or a neurological event being a trigger for a certain stroke.

2.4.1 Relationship with the Cox process

It is worth placing the Hawkes process within a progression of increasingly flexible intensity models. In the homogeneous Poisson process, the intensity λ is constant. In the non-homogeneous Poisson process, it now varies in time in a deterministic way, $\lambda(t)$. One step further is the Cox process, also called the doubly stochastic Poisson process, in which the intensity $\lambda(t)$ itself is a random process, governed by factors external to the event history.

The Hawkes process can be seen as a particular case of this idea of stochastic intensity: the randomness of $\lambda(t)$ does not come from an external source, but from the past events themselves, each one temporarily raising the probability of new events (Equation 2.20). It is in this sense that the Hawkes process is described as doubly stochastic, as mentioned earlier. In this work, it is exactly this dependence on the history (and the exponential decay of the influence of past events) that motivates the recency bias of the proposed model.

2.5 Neural temporal point processes

We saw that Hawkes processes describe many scenarios; however, their classical form does not account for the everyday complexities that are intrinsic to many events. Events may have an inhibitory influence, instead of an excitatory one, on other events, such as, for example, preventive maintenance lowering the probability of a failure in a piece of equipment. Another assumption of traditional Hawkes processes is to consider that events of the same type have a constant influence on other events; however, we know that, for example, when watching a phenomenon for the twentieth time our brain is less excited than when we witness the event for the first time. Finally, we can also recall phenomena where the influence on other events does not decay in an exponential way (MEI; EISNER, 2017).

Mei e Eisner (2017) mention how hard it is to apply the traditional Hawkes process in scenarios with missing data. Even in domains where Hawkes would be appropriate, the nature of the data itself may present challenges, such as partially observed sequences or data omitted in a systematic way. There is also data omitted in a stochastic way, which can be quite complex for the traditional version of Hawkes.

To address such limitations of the Hawkes process, approaches using neural networks were proposed, with the work of Mei e Eisner (2017) being one of the most widespread and accepted by the community, where they proposed the NHP.

The first step was to relax the rule of keeping α and μ strictly positive, and this allowed the values to range over $\mathbb{R}$, addressing inhibition ($\alpha < 0$) and inertia ($\mu < 0$). However, the result can now be negative, which would not make sense for a probability intensity, so it is necessary to apply a transformation $f: \mathbb{R} \rightarrow \mathbb{R}_+$, with $\lambda_k(t) = f_k(\tilde{\lambda}_k(t))$, where k is the event type. The ReLU function $f(x) = \max(x, 0)$ would be a natural choice for this conversion, but it takes the value 0 for negative numbers, and our intensity needs to take always positive values whenever there is a possibility of an event occurring.

The function then chosen by the authors was the softplus $f(x) = \text{slog}(1 + \exp(x/s))$, where s is a learned scale parameter introduced by the model to bring numerical stability, because when the unit of x varies, for example, from seconds to milliseconds, the base intensity $f(\mu_k)$ becomes a thousand times smaller, forcing μ_k to take negative values and thus creating a strong inertia effect.

The second step was to remove the restriction that past events have independence and additive influence on $\lambda(t)$. Instead of predicting $\lambda(t)$ using the summation of the traditional parametric Hawkes model (Equation 2.20), Mei e Eisner (2017) used a recurrent neural network, which allowed them to capture complex patterns that were previously hard to perceive in parametric versions of the model.

As in the traditional models, each event type k has a time-dependent intensity $\lambda_k(t)$, which jumps at each new event and decays exponentially toward a stationary value, analogous to the μ in Equation 2.20. However, in the model proposed by Mei e Eisner (2017) this dynamic is controlled by a hidden state vector $\mathbf{h}(t) \in (-1, 1)^D$, which depends on the vector $\mathbf{c}(t) \in \mathbb{R}^D$, which is the memory vector of a recurrent neural network with one layer and D neurons.

Mei e Eisner (2017) proposed a continuous-time variation of the *Long Short-Term Memory* (LSTM), the Continuous-Time LSTM (CT-LSTM), where after each event each memory cell c decays exponentially at a certain rate δ toward a stable (stationary) state $\bar{\mathbf{c}}$. At each moment $t > 0$, it is possible to obtain the intensity of each event type with $\lambda_k(t) = f_k(\mathbf{w}_k^T \mathbf{h}(t))$, where $\mathbf{h}$ can be computed with $\mathbf{h}(t) = \mathbf{o}_i \odot (2\sigma(2\mathbf{c}(t)) - 1)$ for $t \in (t_{i-1}, t_i]$, where $\mathbf{o}$ is the LSTM output

vector and $\mathbf{w}$ a weight vector.

$$\mathbf{i}_{i+1} \leftarrow \sigma(\mathbf{W}_i \mathbf{k}_i + \mathbf{U}_i \mathbf{h}(t_i) + \mathbf{d}_i) \quad (2.21)$$

$$\mathbf{f}_{i+1} \leftarrow \sigma(\mathbf{W}_f \mathbf{k}_i + \mathbf{U}_f \mathbf{h}(t_i) + \mathbf{d}_f) \quad (2.22)$$

$$\mathbf{z}_{i+1} \leftarrow 2\sigma(\mathbf{W}_z \mathbf{k}_i + \mathbf{U}_z \mathbf{h}(t_i) + \mathbf{d}_z) - 1 \quad (2.23)$$

$$\mathbf{o}_{i+1} \leftarrow \sigma(\mathbf{W}_o \mathbf{k}_i + \mathbf{U}_o \mathbf{h}(t_i) + \mathbf{d}_o) \quad (2.24)$$

$$\mathbf{c}_{i+1} \leftarrow \mathbf{f}_{i+1} \odot \mathbf{c}(t_i) + \mathbf{i}_{i+1} \odot \mathbf{z}_{i+1} \quad (2.25)$$

$$\bar{\mathbf{c}}_{i+1} \leftarrow \bar{\mathbf{f}}_{i+1} \odot \bar{\mathbf{c}}_i + \bar{\mathbf{i}}_{i+1} \odot \bar{\mathbf{z}}_{i+1} \quad (2.26)$$

$$\delta_{i+1} \leftarrow f(\mathbf{W}_d \mathbf{k}_i + \mathbf{U}_d \mathbf{h}(t_i) + \mathbf{d}_d) \quad (2.27)$$

Equations 2.21, 2.22, and 2.24 refer respectively to the *input*, *forget*, and *output gates*, common to a conventional LSTM network. Equation 2.23 refers to the candidate value (*Candidate Memory*). Equations 2.25 and 2.26 represent the Memory Cell (*Memory Cell*) and the stable state of the Memory Cell. Finally, Equation 2.27 computes the vector δ , the decay rate, analogous to the β in Equation 2.20.

The equations above are quite similar to a conventional (discrete) LSTM, but with two fundamental differences. The updates do not depend on the previous value t_{i-1} , but on the value of $\mathbf{h}(t_i)$, after it has undergone decay over the period $t_i - t_{i-1}$. In addition, Equations 2.26 and 2.27 are new, added by Mei e Eisner (2017), and they define how, over time, as $t > t_i$ increases, the elements of $\mathbf{c}(t)$ will continuously decay from $\mathbf{c}_{i+1}$ to $\bar{\mathbf{c}}_{i+1}$. $\mathbf{c}(t)$ is computed according to the equation

$$\mathbf{c}(t) \stackrel{\text{def}}{=} \bar{\mathbf{c}}_{i+1} + (\mathbf{c}_{i+1} - \bar{\mathbf{c}}_{i+1}) \exp(-\delta_{i+1}(t - t_i)) \quad \text{for } t \in (t_i, t_{i+1}]. \quad (2.28)$$

It is important to remember that $\mathbf{c}(t)$ is used in the function $\mathbf{n}(t)$, which is analogous to the LSTM hidden state vector $\mathbf{h}_i$, which in turn is used in the computation of $\lambda_k(t)$, for a given event type k .

The approach of Mei e Eisner (2017) was simple but quite clever, achieving excellent results on both synthetic and real datasets, often still being the best benchmark, despite being a 2017 algorithm.

2.6 Transformers and the attention mechanism

Approaches based on recurrent neural networks (RNN), such as the Continuous-Time LSTM (CT-LSTM) proposed by Mei e Eisner (2017), brought great advances in the modeling

of complex patterns; however, they kept the weak point of RNNs, which is the fact that their processing is sequential. RNNs process one event after another, that is, the computation of the next event (t_{i+1}) depends on the current event (t_i), which makes training the network slow. In some scenarios, with very long sequences, it becomes infeasible to use an RNN. More recent works adapted the Transformer architecture proposed by Vaswani *et al.* (2017), adapting the attention mechanism for TPP, especially aimed at the Hawkes process, which is the focus of this work. The following subsections will help us recall how the architecture proposed in 2017 works.

2.6.1 Encoders and Decoders

The Encoder is the component responsible for processing the input sequence and transforming it into a representation richer in information. It is composed of N identical stacked layers (in the original model of Vaswani *et al.* (2017), $N = 6$ and the dimension of the vectors is $d_{model} = 512$, but these values are hyperparameters that vary according to the application). Each layer has two sublayers: the first is the Multi-Head Self Attention, which computes the relevance of each input with respect to the other inputs in the sequence. The second is a fully connected feed-forward network, which applies nonlinearity at each position.

Similarly, the Decoder is also made up of N stacked layers. However, to generate the output, the Decoder uses a third sublayer, which performs a Multi-Head Attention over the Encoder output, connecting the context processed by the Encoder with the sequence generated by the Decoder. In addition, the Decoder uses a mask to prevent the model from computing future positions in the sequence, so that each prediction considers only known events (Vaswani *et al.*, 2017).

2.6.2 Attention mechanism

The main innovation introduced by the model is the attention mechanism, also called Scaled Dot-Product Attention, without the need for recurrent neural networks, which makes all the inputs (events) be processed simultaneously. In this way the model is able to compute the influence (attention) of each event with respect to the other events.

Unlike RNNs, which keep only a single hidden state vector $\mathbf{h}$, Transformer-based networks compute the influence of each event and compute $\mathbf{h}$ through a weighted sum, which in turn was computed based on the influence of each event, within the context of the list of events.

The input consists of query vectors Q and key vectors K of dimension d_k and value vectors V of dimension d_v . The product of the query with all the keys is then computed, each divided by $\sqrt{d_k}$, and then a softmax function is applied to obtain the weights of the values (of the key-value vectors). The output matrix is computed as

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V. \quad (2.29)$$

2.6.3 Multi-Head Attention

Instead of performing a single Attention function with Queries, Keys, and Values of dimension d_{model} , Vaswani *et al.* (2017) decided to linearly project the Q , K , and V vectors a number h of times with different, learned linear projections to d_k , d_k , and d_v dimensions, respectively. For each of the projected versions of Q , K , and V the Attention function is run in parallel, producing outputs of d_v dimensions. The outputs are then concatenated and projected again, producing the final values, as shown in the equation below:


$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \quad (2.30)$$

where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$


Figure 3 illustrates the difference between the simple attention mechanism and the multi-head attention mechanism (VASWANI *et al.*, 2017).

Figure 3 – Attention x Multi-Head Attention.

Scaled Dot-Product Attention



Multi-Head Attention



3 RELATED WORK

This chapter reviews the Transformer-based point process models that precede and provide the basis for the HoTHP, along with the length extrapolation techniques that motivate its proposal. Here we start from the Transformer Hawkes Process (THP) and the RoTHP, which replaces the absolute temporal *encoding* with *rotary position embeddings*. Finally, we discuss the problem of *length extrapolation* and the notion of a recency inductive bias, which connects these models to the Hawkes kernel and gives rise to the approach proposed in this dissertation.

3.1 Transformer Hawkes Process (THP)

The Transformer architecture (VASWANI *et al.*, 2017) showed that using Attention can offer good advances over the use of recurrent neural networks; however, like traditional RNNs, Transformers were also designed for discrete events, such as the words of a sentence. Zuo *et al.* (2021) had to modify the architecture, turning it into a continuous one, just as Mei e Eisner (2017) had to do when turning the traditional LSTM into a Continuous Time-LSTM.

In translation problems, which were originally the first application of Transformers, the order of the words is enough, with no need for precision beyond that. However, in the analysis of point processes we need to know the moment at which each event occurs. To solve this problem, Zuo *et al.* (2021) proposed a temporal encoding, represented by Equation 3.1, where the argument is not the (discrete) position, as proposed by Vaswani *et al.* (2017), but the timestamp t_j :

$$[z(t_j)]_i = \begin{cases} \cos\left(\frac{t_j}{10000^{\frac{i-1}{M}}}\right), & \text{if } i \text{ is odd,} \\ \sin\left(\frac{t_j}{10000^{\frac{i}{M}}}\right), & \text{if } i \text{ is even.} \end{cases} \quad (3.1)$$

Equation 3.1 uses trigonometric functions to determine a temporal encoding for each timestamp, that is, for each t_j a deterministic $\mathbf{z}(t_j) \in \mathbb{R}^M$ is computed, where M is the size of the encoding dimension. Besides the temporal encoding, an embedding layer, the matrix $\mathbf{U}$, is also used for the different event types. The embeddings are then represented by $\mathbf{X} = (\mathbf{U}\mathbf{Y} + \mathbf{Z})^T$, where $\mathbf{Y}$ is the one-hot encoding collection of event types and $\mathbf{Z}$ the matrix with the $\mathbf{z}$ vectors concatenated. After the embedding layers, $\mathbf{X}$ is passed through a self-attention mechanism as shown below:

$$\mathbf{S} = \text{Softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V}, \quad \text{where } \mathbf{Q} = \mathbf{X}\mathbf{W}^Q, \mathbf{K} = \mathbf{X}\mathbf{W}^K, \mathbf{V} = \mathbf{X}\mathbf{W}^V. \quad (3.2)$$

$\mathbf{Q}$, $\mathbf{K}$, and $\mathbf{V}$ are the Query, Key, and Value matrices obtained through transformations on $\mathbf{X}$, and $\mathbf{W}^Q$, $\mathbf{W}^K$, and $\mathbf{W}^V$ are the weights of linear transformations. Then, $\mathbf{S}$ passes through a feed forward network producing $\mathbf{h}(t)$:

$$\mathbf{H} = \text{ReLU}(\mathbf{S}\mathbf{W}_1^{FC} + \mathbf{b}_1)\mathbf{W}_2^{FC} + \mathbf{b}_2, \quad \mathbf{h}(t_j) = \mathbf{H}(j, :). \quad (3.3)$$

Figure 4 – Architecture of the Transformer Hawkes Process




Figure 4 shows the architecture of the Transformer Hawkes Process. Each event sequence passes through two embedding layers and several times through multi-head self attention modules.

Since it models a multivariate process, the total intensity at a time t is the sum of the intensities of all K event types:

$$\lambda(t | \mathcal{H}_t) = \sum_{k=1}^K \lambda_k(t | \mathcal{H}_t) \quad (3.4)$$

As we saw earlier, after the whole process of embedding, attention, and feed-forward network, $\mathbf{H}$ is finally produced, which lets us compute λ_k , needed in Equation 3.4. The intensity computation for each event type is given by:

$$\lambda_k(t | \mathcal{H}_t) = f_k \left(\underbrace{\alpha_k \frac{t - t_j}{t_j}}_{\text{current}} + \underbrace{\mathbf{w}_k^\top \mathbf{h}(t_j)}_{\text{history}} + \underbrace{b_k}_{\text{inertia}} \right) \quad (3.5)$$

In Equation 3.5, the first "current" part is an interpolation between t_j and t_{j+1} where α , which is a learned element, modulates the importance of this interpolation. The second part

uses the history $\mathbf{h}(t)$ of the events, and the last element β_k represents the stationary state, where the probability of the next event is not influenced by the event history.

Zhang *et al.* (2020), with the *Self-Attentive Hawkes Process* (SAHP), proposed a different computation for the temporal encoding. Unlike the THP, which uses only the timestamp t_i , the SAHP shifts the event to a new position i' , using both the position i and the timestamp t_i , which is computed with $\sin(\omega_k \times i + w_k \times t_i)$, where ω is the angular frequency and w is a scale parameter, adjusted by the model, that converts the timestamp t_i into a phase shift inside the sinusoidal function (as in other models, the Self-Attentive Hawkes Process uses sine for even events and cosine for odd events).

3.2 Rotary Position Embeddings (RoPE) and the RoTHP

We saw that the use of RNNs has limitations both for the memory of distant events and in the training of the network itself, because each event needs to be processed before the next event, since it carries the state of the previous event. Through the attention mechanism, models using Transformers solve these two problems; however, most implementations are not resilient to noise in the data (timestamps). In real datasets, the presence of such noise is very common, mainly due to physical delays in recording the timestamp. For example, in wireless networks, devices record events even without having their clocks perfectly synchronized with one another, which causes discrepancies in the final dataset. The THP and its most common variants adopt a sinusoidal temporal encoding, ignoring such characteristics of the data.


Su *et al.* (2023) showed that the likelihood computation uses only the difference between the event times and, in this way, proposed to use the relative distances between tokens in language models, replacing the temporal encoding of the traditional Transformer architecture with the relative time, inside the attention mechanism. To do this, we use the input vector (embedding) in pairs (d_i, d_{j+1}) . Each pair passes through a 2D rotation matrix, producing a rotation matrix $\mathbf{R}$, where $\Theta = \{\theta_j = 10000^{-2(i-1)/d}, i \in [1, 2, \dots, d/2]\}$, with $f_{\{q,k\}}(\mathbf{x}_m, m) = \mathbf{R}_{\theta, m}^d \mathbf{W}_{\{q,k\}} \mathbf{x}_m$, where

q and k are the Query and Key components of the attention mechanism:

$$R_{\Theta,m}^d = \begin{pmatrix} \cos(t_i\theta_1) & \sin(t_i\theta_1) & 0 & 0 & \cdots & 0 & 0 \\ -\sin(t_i\theta_1) & \cos(t_i\theta_1) & 0 & 0 & \cdots & 0 & 0 \\ 0 & 0 & \cos(t_i\theta_2) & \sin(t_i\theta_2) & \cdots & 0 & 0 \\ 0 & 0 & -\sin(t_i\theta_2) & \cos(t_i\theta_2) & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \cdots & \cos(t_i\theta_{d/2}) & \sin(t_i\theta_{d/2}) \\ 0 & 0 & 0 & 0 & \cdots & -\sin(t_i\theta_{d/2}) & \cos(t_i\theta_{d/2}) \end{pmatrix}.$$

Figure 5 illustrates the tokens of a language model, the position vector, and how RoPE uses pairs of the embedding to apply rotation. The result of the matrix R (rotation) is then applied to the matrices Q and K of the attention mechanism.

Figura 5 – Rotary Position Embeddings (RoPE).



Dai *et al.* (2026) adapted RoPE to point processes, presenting the RoTHP. This consisted of no longer using the index position of the token of a language model, but the timestamp. A great benefit of this approach was making the likelihood time-invariant, since it deals only with the relative times (deltas). This means that if we have a time series $\{t_1 + \sigma, t_2 + \sigma, \dots, t_n + \sigma\}$ it will produce the same result as $\{t_1, t_2, \dots, t_n\}$, and this solves a consistency problem inherent to models that use absolute times, allowing a time shift.

The metrics presented in Tables 1, 2, and 3 show a clear superiority of the RoTHP over several traditional models on several datasets. Although the accuracy is competitive with

the THP, the advance was significant in the log-likelihood computation (NLL) and in the RMSE. In the RMSE the improvement was substantial, especially on the Stack Overflow (SO) datasets, where the RoTHP surpassed the THP, which until then was the best model, by a significant margin. This suggests that the proposed embedding rotation model captures patterns more precisely than the traditional sinusoidal encoding.

In addition, part of the robustness of the RoTHP comes from its ability to remain time-invariant. By replacing absolute timestamps with time intervals through rotation in the 2D plane, the model effectively managed to mitigate the impact of noise, often caused by synchronization problems present in real datasets. This shows that adapting RoPE to point processes increased the generalization capability of current models that use Transformers. As a result, the RoTHP establishes a new state of the art for point process benchmarks, as far as we know.

Tabela 1 – Model log-likelihood by dataset

Models	Financial	SO	Synthetic	Retweet	MemeTrick	Mimic-II
RMTTP	-3.89	-2.6	-1.33	-5.99	-6.04	-1.35
NHP	-3.6	-2.55	–	-5.6	-6.23	-1.38
SAHP	–	-1.86	0.59	-4.56	–	-0.52
THP	-1.11	-0.039	0.791	-2.04	0.68	0.48
RoTHP	1.076	0.389	1.01	2.01	1.71	0.64

Tabela 2 – Accuracy Comparison

Models	Financial	Mimic-II	SO
RMTTP	61.95	81.2	45.9
NHP	62.20	83.2	46.3
THP	62.23	84.9	46.4
RoTHP	62.26	85.5	46.33

Tabela 3 – RMSE Comparison

Models	Financial	Mimic-II	SO
RMTTP	1.56	6.12	9.78
NHP	1.56	6.13	9.83
SAHP	–	3.89	5.57
THP	0.93	0.82	4.99
RoTHP	0.60	0.57	1.33




Figure 6 – As the extrapolation increases, we clearly see that the degradation shown by RoPE is smaller than that of traditional sinusoidal methods, but still much sharper than the model proposed by Press *et al.* (2021)

3.3 Length Extrapolation and inductive bias

A common limitation shared by many Transformer-based sequence models is their inability to generalize to sequences longer than those used during training. This problem is known as *length extrapolation*. Given a model trained on sequences of size L_{train} , we check whether it remains stable when tested with sequences of size $L_{test} > L_{train}$ (PRESS *et al.*, 2021). The practical relevance of this limitation is significant. In many real-world scenarios, collecting long sequences is expensive or even impractical and, therefore, it is desirable to train on shorter sequences and run the model on arbitrarily longer sequences.

The main cause of this failure has been attributed to the *positional encoding* strategy applied in traditional Transformer models. The standard sinusoidal embedding, as introduced by Vaswani *et al.* (2017), produces absolute representations of the positions. When the model encounters a position index that was never presented during training, these representations are effectively considered *out of distribution*, causing a rapid degradation in performance, as shown in the Results section.

Press *et al.* (2021) showed empirically that sinusoidal models are unable to extrapolate to sequences larger than the ones used in training, with the perplexity metric increasing dramatically as L_{test} grows, presenting a model called ALiBi. Rotational embeddings, such as RoPE (SU *et al.*, 2023), present improvements over the sinusoidal encoding by using relative encoding instead of absolute positions, however, Press *et al.* (2021) showed that even RoPE fails to extrapolate to substantially larger sequences, as shown in Figure 6.

3.3.1 ²¹ Attention with Linear Biases (ALiBi)

Press *et al.* (2021) proposed a fundamentally different approach to positional encoding. Instead of adding or rotating embeddings, they introduced a static, non-learnable penalty directly into the attention score. This method, called *Attention with Linear Biases (ALiBi)*, replaces the positional embedding step completely. The main motivation of the model is that it should give less importance to tokens as they move away from the current index, which the authors call *inductive bias towards recency*. Instead of learning this property from the data, ALiBi computes this decay of "importance" in its own architecture.

In traditional Transformer models, ¹⁰⁰ position embeddings are added to the embeddings of the words (tokens). ¹ The dot product of the Q and K vectors is computed as follows:

$$\text{softmax}(\mathbf{q}_i \mathbf{K}^\top), \quad (3.6)$$

where the result is then multiplied by the values of V . ALiBi only modifies the step after the dot product, leaving everything else intact:

$$\text{softmax}(\mathbf{q}_i \mathbf{K}^\top + m \cdot [-(i-1), \dots, -2, -1, 0]), \quad (3.7)$$

where the vector $[-(i-1), \dots, -2, -1, 0]$ represents the distances of each position K to the current Q position i , and the scalar m is a slope coefficient specific to each attention head, fixed before training and never updated. The negative sign guarantees that the penalty lowers the attention score as the distance between Q and K increases, effectively suppressing the attention to distant tokens. For a model with n attention heads, the slope m forms a geometric progression with ratio and first term $2^{-\frac{8}{n}}$ (PRESS *et al.*, 2021).

3.3.2 Inductive bias and its relevance to the Hawkes process

An inductive bias is an assumption that guides the model, independently of the trained data. In the positional context of ALiBi, we can state that the inductive bias dictates that distant words/tokens are less relevant, linearly, than closer ones. The penalty is applied directly in the code, and not only in the learned data.

The inductive bias is not arbitrary. In fact, it has been present in the Hawkes process since its proposal, in 1971. Looking at the Hawkes process equation, the influence of a past event t_i decays as the interval $t - t_i$ increases, in proportion to β . The greater the distance, the

smaller the contribution of the event to the process. This shows that the Hawkes process has at its core an inductive bias represented by the exponential decay parameter.

The THP and the RoTHP replace the parametric Hawkes kernel with an attention mechanism. This brought enormous advantages to these models, such as the possibility of capturing more complex relationships, and even temporal invariance, in the case of the RoTHP, but the price paid was moving away from the formal equation of the original Hawkes process.

The attention mechanism does not impose any constraint on how it should consider the distance between the tokens, in the case of a TPP, between the events. Any decay that emerges from the model is the result of the learning performed with the data and not a structural guarantee. In fact, due to the oscillatory nature of the sinusoidal computation, occasionally distant events may receive more relevance than close events.

4 METHODOLOGY

The HoTHP (*Hyperbolic Rotary Position Embedding-based Transformer Hawkes Process*) is a variant of the THP that replaces the trigonometric kernel of the RoTHP with a hyperbolic kernel that has exponential decay, inspired by the HoPE (DAI *et al.*, 2025). The main motivation is length extrapolation, where we train ¹¹¹on short sequences and keep performance stable on longer sequences. The mechanism behind this ability is an exponential recency bias embedded directly in the attention kernel, which makes the scores decay monotonically with the temporal distance, reflecting the structure of the Hawkes process itself.

4.1 The oscillatory limitation of the RoTHP

As we saw in Section 3.2, the RoTHP represents the position of each event by a trigonometric ⁸³rotation applied to the query $\mathbf{q}$ and key $\mathbf{k}$ vectors. The attention score between events i and j depends only on $\Delta t = t_i - t_j$, which is guaranteed by the algebraic properties of the trigonometric rotation matrix, which is orthogonal.

The problem is that the resulting kernel is oscillatory. Sine and cosine are not monotonic functions. As a result, the attention score can grow for events more distant in time, depending on the phase of the oscillation, and shrink for nearby events. This directly contradicts the kernel ¹¹³of the Hawkes process, $\varphi(\Delta t) = \alpha e^{-\beta \Delta t}$, which is strictly decreasing.

In practice, the most serious consequence of this oscillation appears in length extrapolation. When tested on sequences longer than the training ones, the model encounters temporal differences Δt that it never observed during training. A kernel with exponential decay would have a predictable behavior in this region, with the influence continuing to fall. The trigonometric kernel, on the other hand, keeps oscillating with the same amplitude, and a very distant event can receive a high score simply because the phase of the cosine fell on a peak at that Δt . The model did not learn to handle these phases because it never saw them during training, and the kernel has nothing that forces the score to decay. It is this absence of a recency bias that prevents the RoTHP from extrapolating in a stable way.

4.2 From the trigonometric kernel to the hyperbolic kernel

For each pair of dimensions $(2j-1, 2j)$, the RoTHP applies the 2×2 trigonometric rotation block:

$$\mathbf{R}_j(\Delta t) = \begin{bmatrix} \cos(\theta_j \Delta t) & \sin(\theta_j \Delta t) \\ -\sin(\theta_j \Delta t) & \cos(\theta_j \Delta t) \end{bmatrix}. \quad (4.1)$$


The HoPE (DAI *et al.*, 2025) proposes to replace this rotation with a hyperbolic rotation, whose corresponding block is:

$$\mathbf{B}_j(\Delta t) = \begin{bmatrix} \cosh(\theta_j \Delta t) & \sinh(\theta_j \Delta t) \\ \sinh(\theta_j \Delta t) & \cosh(\theta_j \Delta t) \end{bmatrix}. \quad (4.2)$$

Unlike $\cos$ and $\sin$, the functions $\cosh$ and $\sinh$ are monotonic for $\Delta t > 0$, removing the oscillations. However, we cannot use $\mathbf{B}_j$ directly because the hyperbolic rotation is not orthogonal and the inner product $\mathbf{q}^\top \mathbf{B}(\Delta t) \mathbf{k}$ grows with $|\Delta t|$ instead of decaying (DAI *et al.*, 2025). Distant events would receive larger scores than nearby events, exactly the opposite of what we want.

Figure 7 illustrates this difference. In the left panel, $\cos$ and $\sin$ oscillate between -1 and 1 indefinitely, without any decay trend. In the right panel, $\cosh$ and $\sinh$ remove the oscillations but grow exponentially in both directions. Note that $\cosh$ is an even function (symmetric, with minimum value $\cosh(0) = 1$) while $\sinh$ is odd (antisymmetric, with $\sinh(0) = 0$).

Figura 7 – Base functions of the attention kernels ($\theta = 1, 0$). Left: trigonometric functions of the RoTHP, which oscillate between -1 and 1 without decaying. Right: hyperbolic functions of the HoPE, which are monotonic but grow exponentially. The growth of the hyperbolic functions needs to be compensated by a decay factor so that the attention kernel works correctly.



4.3 Exponential decay and the parameter θ'

To correct the hyperbolic growth, the HoPE (DAI *et al.*, 2025) introduces a learnable parameter θ' and applies damping factors to the $\mathbf{q}$ and $\mathbf{k}$ vectors. The vector $\mathbf{q}$ at position m is multiplied by $e^{-m\theta'}$ and the vector $\mathbf{k}$ at position n is multiplied by $e^{+n\theta'}$. When computing the score between positions m and n , these factors combine:

$$g(\mathbf{x}_m, \mathbf{x}_n, n-m) = e^{-(m-n)\theta'} \mathbf{q}^\top \mathbf{B}(\theta, m-n) \mathbf{k}. \quad (4.3)$$

In the HoTHP, the discrete positions m, n are replaced by the continuous instants t_i, t_j and the difference $m - n$ by the temporal difference $\Delta t = t_i - t_j$:

$$g(\Delta t) = e^{-\Delta t \cdot \theta'} \mathbf{q}^\top \mathbf{B}(\theta, \Delta t) \mathbf{k}. \quad (4.4)$$

For the decay to dominate the hyperbolic growth, it is necessary that $\theta' > \max_j \theta_j$. Under this condition, the score decreases monotonically with $|\Delta t|$, which introduces an exponential recency bias into the attention mechanism. This bias reflects the same structure as the Hawkes kernel ($\alpha e^{-\beta \Delta t}$) and, because it is structural and not learned from the data, it holds even for values of Δt that the model never saw during training. It is this bias that gives the HoTHP the ability to extrapolate to longer sequences without degradation.

Figure 8 illustrates the recency bias in practice. The left image shows the separate components: cosh and sinh grow exponentially while the decay factor $e^{-\theta' \Delta t}$ falls quickly to zero. The center image shows the product of the hyperbolic functions by the decay factor, ensuring that $\theta' > \theta$, the decay dominates the growth, and both $e^{-\theta' \Delta t} \cosh(\theta \Delta t)$ and $e^{-\theta' \Delta t} \sinh(\theta \Delta t)$ decay monotonically. The right image shows the decomposition into pure exponentials used in the implementation (Section 4.5): since both exponents $(\theta' - \theta)$ and $(\theta' + \theta)$ are positive, all the exponentials have a negative argument and no overflow occurs.

4.3.1 Parametrization of the learnable θ'


Since θ' needs to be learnable but at the same time satisfy $\theta' > \max_j \theta_j$, we cannot simply train a free parameter. The frequencies θ_j follow the RoPE formula (SU *et al.*, 2023):

$$\theta_j = 10000^{-2(j-1)/d}, \quad j = 1, \dots, d/2, \quad (4.5)$$

where d is the dimension of each attention head. Each pair of dimensions receives a different frequency, covering multiple scales of Δt . Note that, by Equation 4.5, the largest value of θ_j

Figura 8 – Effect of the exponential decay on the hyperbolic functions ($\theta = 1,0$, $\theta' = 1,5$). Left: separate components, showing that cosh and sinh grow but the decay falls. Center: the product results in functions that decay monotonically. Right: decomposition into pure exponentials, both with a negative exponent, ensuring numerical stability.

HoTHP: Decaimento Exponencial Domina o Crescimento Hiperbolico
($\theta = 1,0$, $\theta' = 1,5$, com $\theta' > \theta$)



occurs when $j = 1$, resulting in $\theta_1 = 10000^0 = 1$. All the other θ_j are increasingly smaller fractions (for example, $\theta_2 \approx 0,013$, $\theta_3 \approx 0,00017$, and so on). Therefore, in practice, the constraint $\theta' > \max_j \theta_j$ reduces to ensuring that $\theta' > 1$. Even so, we chose a general formulation that works regardless of the values of θ_j , so that the constraint is satisfied by construction. The solution is to parametrize θ' as:

$$\theta' = \text{softplus}(\theta'_{\text{raw}}) + \max_j(\theta_j) + \varepsilon, \quad (4.6)$$

where θ'_{raw} is the trainable parameter, $\text{softplus}(x) = \ln(1 + e^x)$ ensures that the first term is positive, and $\varepsilon = 10^{-4}$ is a safety margin. Thus, θ' is always larger than any θ_j and the exponential decay is structurally guaranteed, regardless of what the gradient does.

4.4 Attention score of the HoTHP

Expanding the matrix multiplication of $\mathbf{B}_j$ (Equation 4.2) over the even-dimension vectors $\mathbf{q}_1, \mathbf{k}_1$ and the odd-dimension vectors $\mathbf{q}_2, \mathbf{k}_2$, we obtain:

$$\mathbf{q}^\top \mathbf{B}_j(\Delta t) \mathbf{k} = \underbrace{(\mathbf{q}_1 \cdot \mathbf{k}_1 + \mathbf{q}_2 \cdot \mathbf{k}_2)}_{\text{cosh part}} \cdot \cosh(\theta_j \Delta t) + \underbrace{(\mathbf{q}_1 \cdot \mathbf{k}_2 + \mathbf{q}_2 \cdot \mathbf{k}_1)}_{\text{sinh part}} \cdot \sinh(\theta_j \Delta t). \quad (4.7)$$

The complete score incorporates the decay factor and sums over the $d/2$ pairs of dimensions:

$$s_{ij} = \frac{1}{\sqrt{d}} \sum_{j=1}^{d/2} \left[(q_1^{(j)} k_1^{(j)} + q_2^{(j)} k_2^{(j)}) \cdot e^{-|\Delta t| \theta'} \cosh(\theta_j \Delta t) + (q_1^{(j)} k_2^{(j)} + q_2^{(j)} k_1^{(j)}) \cdot e^{-|\Delta t| \theta'} \sinh(\theta_j \Delta t) \right], \quad (4.8)$$

where $\Delta t = t_i - t_j$ and $1/\sqrt{d}$ is the standard normalization of the **scaled dot-product attention** (VASWANI *et al.*, 2017).

In the RoTHP, the $\mathbf{q}$ and $\mathbf{k}$ vectors are pre-rotated individually by $\mathbf{R}(t_i)$ and $\mathbf{R}(t_j)$ before the inner product, and the dependence on Δt emerges from the algebraic cancellation demonstrated by the authors. In the HoTHP, the score is formulated directly in terms of Δt , which makes it natural and necessary to include the factor $e^{-|\Delta t|\theta'}$ inside the kernel itself. The practical consequence is that, for any Δt (including values larger than those seen in training), the score is dominated by the exponential decay. This is the mechanism that enables extrapolation: the recency bias does not depend on the data, it is in the structure of the attention.

4.5 Numerical stability

The functions $\cosh$ and $\sinh$ grow like $e^{|\Delta t|\theta_j}$ and, for events that are very far apart, these values can cause *overflow* even before being multiplied by the decay factor $e^{-|\Delta t|\theta'}$.

The solution is to distribute the decay into the hyperbolic functions using the identities:

$$e^{-|\Delta t|\theta'} \cdot \cosh(|\Delta t|\theta_j) = \frac{e^{-|\Delta t|(\theta' - \theta_j)} + e^{-|\Delta t|(\theta' + \theta_j)}}{2}, \quad (4.9)$$

$$e^{-|\Delta t|\theta'} \cdot \sinh(\Delta t \cdot \theta_j) = \operatorname{sgn}(\Delta t) \cdot \frac{e^{-|\Delta t|(\theta' - \theta_j)} - e^{-|\Delta t|(\theta' + \theta_j)}}{2}. \quad (4.10)$$

Since $\theta' > \max_j \theta_j$, we have $\theta' - \theta_j > 0$ and $\theta' + \theta_j > 0$, so both exponents are negative for any Δt and the *overflow* does not occur. The sign of Δt is reintroduced explicitly by $\operatorname{sgn}(\Delta t)$, since $\sinh(-x) = -\sinh(x)$.

4.6 Temporal normalization

Different datasets have very distinct temporal scales. In social networks, the intervals between events can be on the order of milliseconds, while in seismic data they are on the order of days. If we used the raw instants in the hyperbolic kernel, the values of $\theta_j \Delta t$ would be on incompatible scales and the model would need specific hyperparameters for each dataset.

To avoid this, we apply two transformations to each sequence. First, we shift the instants so that the sequence starts at zero and, next, we divide all the instants by the mean of the inter-event intervals of that sequence:

$$\tilde{t}_i = \frac{t_i - t_1}{\bar{g}}, \quad (4.11)$$

where $\bar{g}$ is the arithmetic mean of the intervals between consecutive events of the sequence. This operation is done independently for each sequence. The shift ensures that $\tilde{t}_1 = 0$ and the

division by $\bar{g}$ makes the mean interval between events become approximately 1 in all sequences, regardless of the original temporal scale of the dataset. The order and the proportion between the intervals are preserved.

The choice of this normalization was made after we evaluated two alternatives. The first was a MinMax normalization applied to the *intervals* between events, and not to the absolute instants: each $\Delta t_i = t_i - t_{i-1}$ is mapped to the interval $[0, 1]$ by $\Delta \tilde{t}_i = (\Delta t_i - \Delta t_{\min}) / (\Delta t_{\max} - \Delta t_{\min})$, where $\Delta t_{\min}$ and $\Delta t_{\max}$ are the smallest and the largest interval of the sequence. Applying the MinMax directly to the absolute instants would be inadequate, because the last event is always the maximum and, as the sequence grows, all the previous instants are pushed close to zero; normalizing the intervals avoids this problem. Even so, this alternative proved inadequate for the *train-short, test-long* protocol: in the synthetic extrapolation experiments it systematically favored the NHP, which does not depend on explicit positional embeddings, while the models with temporal attention (THP, RoTHP, and HoTHP) were harmed or showed instability, especially at the high extrapolation factors.

The reason is that the MinMax divisor depends on the *maximum* interval of the sequence, and the maximum is an extreme value that tends to grow the more events the sequence has, since longer sequences have a greater chance of containing an atypically large interval. The mean normalization does not suffer from this problem because the divisor $\bar{g}$ is the mean of the intervals, a stable statistic that barely changes when more events from the same process are added: the normalized mean interval stays ≈ 1 at any length. An atypical interval is also diluted by the mean, whereas in the MinMax it would change the maximum and compress all the others.

The second alternative evaluated was a temporal scale factor learned as a parameter of the model during training. This approach introduced instability in the optimization, with high variance between seeds at the high extrapolation factors. The mean-interval normalization, because it is simple and does not depend on the length of the test sequence, was adopted in the experiments with synthetic data and in the Retweet ablation (Section 5.4).

4.7 Complete architecture

The HoTHP architecture follows the Transformer encoder (VASWANI *et al.*, 2017). Each event (t_i, k_i) is represented by a learnable type embedding $\mathbf{e}_{k_i} \in \mathbb{R}^d$. Unlike the THP (ZUO *et al.*, 2021), which adds a sinusoidal encoding of the absolute time to the embedding, the HoTHP does not include any temporal information in the embedding. The instants t_i enter the

model only through the hyperbolic kernel.

The model stacks N hyperbolic encoder layers, each one with a hyperbolic multi-head attention block followed by a *feed-forward* network. The attention computes the scores via Equation 4.8, using the normalized differences $\Delta \tilde{t}_{ij} = \tilde{t}_i - \tilde{t}_j$. A causal mask ensures that each event attends only to previous events.

The *feed-forward* network in each layer has two linear projections with a ReLU between them:

$$\text{FFN}(\mathbf{x}) = \mathbf{W}_2 \cdot \text{ReLU}(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2, \quad (4.12)$$

with internal dimension $2d$. Like the RoTHP (DAI *et al.*, 2026), the HoTHP does not use residual connections (add & norm).

The encoder output $\mathbf{h}(t_j) \in \mathbb{R}^d$ for event j is used to compute the conditional intensity of each type k at the instant $t > t_j$:

$$\lambda_k(t | \mathcal{H}_t) = \text{softplus} \left(\alpha_k \frac{t - t_j}{t_j + \varepsilon} + \mathbf{w}_k^\top \mathbf{h}(t_j) + b_k \right), \quad (4.13)$$

where α_k , $\mathbf{w}_k$, and b_k are learnable parameters and softplus ensures that the intensity is positive. The term $\alpha_k(t - t_j)/t_j$ captures the variation with the time since the last event, and $\mathbf{w}_k^\top \mathbf{h}(t_j)$ brings the historical context from the encoder. This form follows the THP (ZUO *et al.*, 2021).

4.8 Where the exponential bias acts

An important distinction needs to be made about the role of the exponential decay in the HoTHP. The factor $e^{-|\Delta t| \theta'}$ does not act directly on the intensity function $\lambda_k(t)$, but on the attention mechanism that produces the hidden state $\mathbf{h}(t_j)$. These are two distinct steps of the model, with different roles:


1. **Hyperbolic attention** (Equation 4.8): determines the weights with which each past event contributes to the hidden state. The exponential decay makes recent events receive larger weights and distant events receive weights close to zero. The result is the vector $\mathbf{h}(t_j)$, which encodes the historical context of the sequence.
2. **Intensity function** (Equation 4.13): uses the hidden state $\mathbf{h}(t_j)$ and the time elapsed since the last event to compute $\lambda_k(t)$. The parameter α_k in this equation is learnable and can take positive or negative values. If $\alpha_k > 0$, the intensity grows with the time since the last event, characterizing a self-inhibiting process. If $\alpha_k < 0$, the intensity decreases, as in a classical self-exciting Hawkes process.

This separation has an important consequence, which is that the exponential bias of the HoTHP does not require the modeled process to be self-exciting. The hyperbolic kernel only requires that the relevance of past events for the construction of the hidden state decay exponentially with the temporal distance. The dynamics of the intensity itself (whether it rises or falls after an event) is determined by the learnable parameters of the intensity layer, which are free to adapt to the data.

Therefore, the HoTHP is suitable for any point process in which the most recent event carries more information than distant events for predicting the next one. This includes both self-exciting processes (where events generate more nearby events) and self-inhibiting ones (where the occurrence of an event temporarily reduces the probability of new events). The bias is harmful when distant events are as relevant as, or more relevant than, recent ones, as in processes without temporal memory or with long-range dependencies.

Figure 9 visually summarizes the difference between the attention kernels of the RoTHP and the HoTHP. In the training zone (green background), both models observe the same temporal distances and can learn appropriate attention weights. In the extrapolation zone (red background), the trigonometric kernel of the RoTHP keeps oscillating in phases that the model never saw during training, assigning potentially high scores to very distant events. The hyperbolic kernel of the HoTHP, on the other hand, decays monotonically to zero regardless of the distance, thanks to the structural bias. For reference, the dashed curve shows the classical Hawkes kernel ($e^{-\beta\Delta t}$), whose shape the HoTHP reproduces by its own architecture.

Figura 9 – Comparison of the attention kernels as a function of the temporal distance Δt . In the training zone, both models can fit the data. In the extrapolation zone, the trigonometric kernel of the RoTHP oscillates in phases that were not learned, while the hyperbolic kernel of the HoTHP decays monotonically, like the classical Hawkes kernel.



4.9 Experimental setup

The experiments use two types of data: synthetic and real. The synthetic data is generated by Hawkes processes with an exponential kernel. The choice of synthetic data is deliberate: since we know the generating process, we can check whether the recency bias of the HoTHP actually aligns with the real decay of the data and whether this correspondence is what sustains the extrapolation. The real data are the datasets provided by the EasyTPP library (XUE *et al.*, 2024): Taxi, StackOverflow, Retweet, and Amazon. These datasets cover varied domains (transportation, question forums, social networks, and electronic commerce) and allow us to evaluate the performance of the HoTHP in scenarios where the generating process is unknown.

4.9.1 Synthetic data generation

We generate sequences from multivariate Hawkes processes with $K = 2$ event types and an exponential kernel. We define two generation scenarios to cover distinct regimes of temporal dependence.

In the **slow decay scenario**, the parameter β is small, so that the influence of past events persists for longer and distant events still affect the present intensity.

In the **fast decay scenario**, β is large and only recent events have relevant influence.

For each scenario, we generate 400 training sequences, 100 validation sequences, and 200 test sequences. All the sequences are preprocessed with the temporal normalization of Equation 4.11 before being fed to the models.

4.9.2 Train short, test long protocol

The *train short, test long* protocol, proposed by Press *et al.* (2021) for language models, trains the model on short sequences and tests it on longer sequences. We adapted the protocol to point processes using the number of events as the measure of length.

We fix the training length at $L = 50$ events and vary the extrapolation factor over $\{1\times, 2\times, 5\times, 10\times\}$, where the factor indicates how many times longer the test sequence is than the training one. With a factor of $10\times$, for example, we train with sequences of 50 events and test with sequences of 500. Each factor is evaluated in the two decay scenarios.

4.9.3 Compared models and implementation

On the synthetic data we compare four models: NHP (MEI; EISNER, 2017), THP (ZUO *et al.*, 2021), RoTHP (DAI *et al.*, 2026), and the proposed HoTHP. The RoTHP is the closest baseline, since it shares exactly the same encoder structure and the same intensity function as the HoTHP, differing only in the positional kernel; the NHP (recurrent) and the THP (absolute temporal encoding) serve to situate the extrapolation behavior with respect to distinct architectures. On the real data we compare the four models (NHP, THP, RoTHP, and HoTHP) on Taxi, Amazon, and Stackoverflow, and the three Transformer models (THP, RoTHP, and HoTHP) on Retweet. All four models compute the NLL with the same likelihood routine (the same Monte Carlo estimate of the compensator integral, defined once in the base class), so the comparison of likelihoods does not depend on the way each model approximates the integral.

All the models are implemented in the EasyTPP framework (XUE *et al.*, 2024), and the baselines share the same hyperparameters as the HoTHP in each scenario, so as to isolate the effect of the positional kernel. In the synthetic experiments we use a model dimension $d = 32$, a temporal *embedding* dimension equal to 32, $N_\ell = 2$ layers, $H = 2$ attention heads, *dropout* of 0.1, and the AdamW optimizer (weight decay 10^{-4}) with learning rate 10^{-3} , except the HoTHP, which uses 5×10^{-4} for greater numerical stability of the hyperbolic functions. In the experiments with real data we use $d = 64$, a temporal *embedding* equal to 16, $N_\ell = 2$ layers, $H = 2$ heads, and a learning rate of 10^{-3} for all the models.

4.9.4 Metrics and statistical evaluation

The evaluation metric is the NLL, computed over the complete test sequence, including the events beyond the training length. If the recency bias of the HoTHP indeed enables extrapolation, the NLL should remain stable even when the test sequence is much longer than the training one. We also use RMSE to evaluate the prediction of the next event and Accuracy to evaluate the prediction of the type of the next event.

We repeat the synthetic experiments with $N = 10$ seeds and the experiments with real data with $N = 3$ seeds, reporting mean $\pm$ standard deviation. To check statistical significance in the synthetic experiments, we use the Wilcoxon signed-rank test, a non-parametric paired test that does not assume normality and is suitable for small samples. The pairing is by seed: the same seed generates the same training, validation, and test set for all the models. With $N = 10$,

the smallest attainable one-sided p -value is $1/2^{10} \approx 0,001$, sufficient for the level $\alpha = 0,05$. 62

4.9.5 *Computational environment*

All the experiments were run on a machine with an NVIDIA A100-SXM4-80GB GPU (80 GB of VRAM), 6 physical cores (12 threads), 167 GB of RAM, and CUDA 12.8. The code uses PyTorch 2.10.0 and Python 3.12.

5 RESULTS

This chapter presents the experimental results of the HoTHP in two contexts: synthetic data, where the generating process is known, and real data, where it is not. Next, we visualize the intensities learned by the models to understand the behavior of the recency bias in practice, and we run an ablation to isolate the effect of the temporal scale.

5.1 Synthetic data


The synthetic experiments follow the protocol described in Section 4.9, using Hawkes processes with an exponential kernel in the slow decay (β small) and fast decay (β large) scenarios. We compare four models: NHP, THP, RoTHP, and HoTHP. Table 4 summarizes the results for the *train short, test long* protocol, varying the extrapolation factor α from $1\times$ (training length) to $10\times$. Figure 10 visualizes the evolution of the NLL as a function of α .

Tabela 4 – NLL (mean $\pm$ standard deviation, 10 seeds) on the synthetic data. Lower values indicate better performance. The best result for each configuration is in bold.

Scenario	Model	$1\times$	$2\times$	$5\times$	$10\times$
Slow decay	NHP	$1,470 \pm 0,007$	$2,197 \pm 0,131$	$2,667 \pm 0,236$	$2,817 \pm 0,273$
	THP	$1,475 \pm 0,007$	$3,926 \pm 1,505$	$4,127 \pm 1,803$	$4,199 \pm 2,117$
	RoTHP	$1,499 \pm 0,014$	$2,007 \pm 0,284$	$2,642 \pm 0,509$	$3,001 \pm 0,625$
	HoTHP	$1,545 \pm 0,025$	$1,609 \pm 0,096$	$1,943 \pm 0,313$	$2,364 \pm 0,625$
Fast decay	NHP	$1,663 \pm 0,005$	$1,833 \pm 0,134$	$2,057 \pm 0,223$	$2,153 \pm 0,250$
	THP	$1,677 \pm 0,011$	$3,391 \pm 0,781$	$3,761 \pm 1,072$	$3,897 \pm 1,509$
	RoTHP	$1,681 \pm 0,007$	$2,039 \pm 0,279$	$2,146 \pm 0,309$	$2,275 \pm 0,372$
	HoTHP	$1,689 \pm 0,008$	$1,701 \pm 0,012$	$1,740 \pm 0,050$	$1,774 \pm 0,092$

With the inclusion of the NHP and the THP as baselines, four distinct behaviors emerge in the extrapolation. The THP is the model that degrades the most: its NLL jumps to $\approx 3,9$ already at $2\times$ and stays between $\approx 3,9$ and $\approx 4,2$ at the larger factors, in both scenarios, with a very high standard deviation (up to $\pm 2,12$) caused by seeds that diverge completely. This confirms that the absolute temporal encoding of the THP produces representations outside the distribution for new positions. The NHP, despite not using attention, is the most competitive baseline in the extrapolation, probably because its recurrent architecture accumulates information incrementally, without depending on absolute positions; even so, it stays behind the HoTHP at all extrapolation factors in both scenarios.

Figura 10 – NLL (mean $\pm$ standard deviation, 10 seeds) as a function of the extrapolation factor α in the synthetic scenarios. The HoTHP keeps the NLL more stable as α grows, while the THP degrades drastically and the RoTHP and NHP show intermediate degradation.



The HoTHP obtains the best NLL at all extrapolation factors ($2\times$, $5\times$, and $10\times$) in both scenarios. The cost of this advantage appears in distribution: at $1\times$, the HoTHP is the model with the worst NLL (for example, 1,545 against 1,470 of the NHP in the slow decay), a small difference that reverses as soon as the test sequence goes beyond the training length. In the fast decay, the NLL of the HoTHP grows by only 0,085 between $1\times$ and $10\times$ ($1,689 \rightarrow 1,774$), staying practically flat, while the second best model (NHP) grows by 0,490 ($1,663 \rightarrow 2,153$). This stability is expected, since the generating process has an exponential kernel, the recency bias of the HoTHP is perfectly aligned with the real structure of the data. For temporal distances not seen in training, the hyperbolic kernel keeps decaying monotonically, while the trigonometric kernel of the RoTHP oscillates in phases that were not learned and the absolute encoding of the THP diverges.


The advantage of the HoTHP over the three baselines is statistically significant (Wilcoxon signed-rank test, paired by seed and one-sided, $p < 0,05$) at all extrapolation factors of both scenarios, without exception. The least comfortable case is the slow decay at $10\times$ against the NHP ($p = 0,042$); in all the other comparisons the p -value is 0,025 or below, reaching the floor of 0,001 (equal to $1/2^{10}$, the smallest possible value for the one-sided test with 10 seeds) in most of them.

In the slow decay, where the influence of distant events persists for longer, the HoTHP still obtains the best mean NLL at all factors, but its degradation is much larger than in the fast decay: the NLL grows by 0,819 between $1\times$ and $10\times$, against only 0,085 in the fast decay. At $10\times$, the mean of the HoTHP (2,364) is lower than that of the NHP (2,817) and the difference remains statistically significant, even if by the narrowest margin of the whole

experiment ($p = 0,042$), reflecting the high variance between seeds in this regime ($\pm 0,625$). This suggests that, when the memory of the process is long, the exponential decay imposed by the hyperbolic kernel may discard events that are still relevant too early, reducing the advantage margin.

Figure 11 shows the RMSE of the next event prediction. In the slow decay, all the models degrade in a similar way ($\text{RMSE} \approx 2,7\text{--}2,8$ at $10\times$), with the HoTHP showing the lowest RMSE (2,714). In the fast decay, the four models keep a stable and close RMSE (between $\approx 1,40$ and $1,55$ at $10\times$): the RoTHP obtains the lowest value (1,407) and the NHP the highest (1,549), the latter with a slightly higher variance due to an atypical seed. Unlike the NLL, the RMSE does not separate the models clearly. This indicates that the advantage of the HoTHP shows up mainly in the likelihood, that is, in modeling the full temporal distribution, and not in the pointwise prediction of the next interval, where the attention models are practically tied.

Figura 11 – RMSE of the next event prediction (mean $\pm$ standard deviation, 10 seeds) as a function of the extrapolation factor α . Unlike the NLL, the RMSE does not separate the models clearly: in the slow decay all of them degrade in a similar way and in the fast decay they stay close to each other.



5.2 Real data

To evaluate the HoTHP in more realistic scenarios, we ran experiments on four public point process datasets. Table 5 describes each dataset.

We follow the same *train short, test long* protocol: we train each model with sequences truncated at L_{train} events and evaluate at the training length ($1\times$) and at two extrapolation factors ($2\times$ and $5\times$). We use three seeds (42, 142, 242), 100 epochs, and the same hyperparameters as in Section 4.9.3. On Taxi, Amazon, and Stackoverflow we compare the four models (THP, RoTHP, HoTHP, and NHP); on Retweet we compare the three Transformer models (THP,

Tabela 5 – Real datasets used in the experiments. L_{train} is the maximum training length (number of events) and $L_{5\times}$ is the test length in the extrapolation condition.

Dataset	Types	Sequences (train)	L_{train}	$L_{5\times}$	Domain
Amazon	16	6.454	18	90	Product reviews
Stackoverflow	22	1.401	20	100	User badges
Taxi	10	1.400	7	38	Taxi rides
Retweet	3	20.000	50	250	Retweet cascades

RoTHP, and HoTHP), since this dataset is used mainly to study numerical stability (Section 5.4).

It is important to note how the NLL is computed. In our implementation, all four models share the same routine for the log-likelihood: the integral of the total intensity (the compensator) is estimated by the same Monte Carlo rule, with the same number of sample points per interval, defined once in the base class. The NHP only changes the way the intensity itself is produced (a continuous-time recurrent network instead of attention). The differences in NLL reported below therefore come from the models, not from different ways of approximating the integral. With only three seeds we report the mean and the standard deviation, but we do not apply the Wilcoxon test on the real data: with $N = 3$ the smallest attainable one-sided p -value is $1/2^3 = 0,125$, above the level $\alpha = 0,05$, so the significance test is reserved for the synthetic experiments with $N = 10$ seeds.

Table 6 presents the NLL and Table 7 the type-prediction accuracy.

Two layers of comparison appear in these tables. The first is the comparison *within the Transformer family* (THP, RoTHP, and HoTHP), which share the same architecture and the same intensity function and differ only in the positional kernel. This is the controlled comparison that isolates the effect of our contribution. The second is the comparison against the NHP, a recurrent model of a different family, included to situate the behavior across architectures.

Within the Transformer family. On Amazon, the HoTHP obtains the best NLL among the three at every factor and, more importantly, it is the only one that keeps improving as the sequence grows: its NLL drops monotonically from 2,263 at $1\times$ to 2,207 at $2\times$ and 2,187 at $5\times$, while the THP and the RoTHP improve at $2\times$ but worsen again at $5\times$, returning to about 2,25. The type accuracy follows the same trend, rising from 0,312 to 0,340, the largest value among all the models. On Taxi, the HoTHP again has the best NLL of the three at every factor and is the most stable in extrapolation: at $5\times$ the RoTHP degrades to $-0,137$ with a very high standard deviation ($\pm 0,148$, a sign that some seeds diverge), while the HoTHP stays at $-0,266 \pm 0,007$. On Stackoverflow the three models are practically tied (differences within the

Tabela 6 – NLL on the real datasets (mean $\pm$ standard deviation, 3 seeds), for the extrapolation factors $1\times$, $2\times$, and $5\times$. Lower is better. The best value in each column is in bold. NHP is a recurrent baseline (CT-LSTM); the other three share the same Transformer architecture and intensity function and differ only in the positional kernel.

Dataset	Model	NLL $1\times$	NLL $2\times$	NLL $5\times$
Taxi	THP	$-0,273 \pm 0,002$	$-0,314 \pm 0,004$	$-0,243 \pm 0,015$
	RoTHP	$-0,271 \pm 0,003$	$-0,317 \pm 0,005$	$-0,137 \pm 0,148$
	HoTHP	$-0,278 \pm 0,001$	$-0,328 \pm 0,004$	$-0,266 \pm 0,007$
	NHP	$-0,449 \pm 0,004$	$-0,506 \pm 0,004$	$-0,460 \pm 0,003$
Amazon	THP	$2,264 \pm 0,003$	$2,237 \pm 0,002$	$2,254 \pm 0,002$
	RoTHP	$2,268 \pm 0,004$	$2,230 \pm 0,003$	$2,247 \pm 0,002$
	HoTHP	$2,263 \pm 0,004$	$2,207 \pm 0,003$	$2,187 \pm 0,002$
	NHP	$0,658 \pm 0,229$	$0,630 \pm 0,227$	$0,638 \pm 0,224$
Stackoverflow	THP	$2,744 \pm 0,029$	$2,642 \pm 0,025$	$2,519 \pm 0,013$
	RoTHP	$2,766 \pm 0,035$	$2,655 \pm 0,026$	$2,545 \pm 0,018$
	HoTHP	$2,758 \pm 0,025$	$2,649 \pm 0,022$	$2,536 \pm 0,021$
	NHP	$2,742 \pm 0,012$	$2,633 \pm 0,013$	$2,524 \pm 0,020$

Tabela 7 – Type-prediction accuracy on the real datasets (mean $\pm$ standard deviation, 3 seeds), for the factors $1\times$, $2\times$, and $5\times$. Higher is better. The best value in each column is in bold.

Dataset	Model	Acc $1\times$	Acc $2\times$	Acc $5\times$
Taxi	THP	$0,900 \pm 0,001$	$0,917 \pm 0,003$	$0,911 \pm 0,002$
	RoTHP	$0,898 \pm 0,000$	$0,916 \pm 0,002$	$0,907 \pm 0,011$
	HoTHP	$0,897 \pm 0,001$	$0,917 \pm 0,003$	$0,915 \pm 0,004$
	NHP	$0,898 \pm 0,002$	$0,919 \pm 0,001$	$0,918 \pm 0,001$
Amazon	THP	$0,308 \pm 0,001$	$0,317 \pm 0,003$	$0,314 \pm 0,006$
	RoTHP	$0,308 \pm 0,002$	$0,323 \pm 0,002$	$0,320 \pm 0,007$
	HoTHP	$0,312 \pm 0,001$	$0,331 \pm 0,001$	$0,340 \pm 0,001$
	NHP	$0,293 \pm 0,004$	$0,306 \pm 0,006$	$0,315 \pm 0,006$
Stackoverflow	THP	$0,473 \pm 0,001$	$0,466 \pm 0,001$	$0,450 \pm 0,001$
	RoTHP	$0,473 \pm 0,004$	$0,464 \pm 0,002$	$0,448 \pm 0,001$
	HoTHP	$0,474 \pm 0,003$	$0,466 \pm 0,002$	$0,451 \pm 0,001$
	NHP	$0,469 \pm 0,003$	$0,462 \pm 0,002$	$0,445 \pm 0,002$

error margin), with the HoTHP slightly ahead in accuracy at every factor; this dataset does not seem to have a temporal structure that separates the positional biases.

Against the NHP. The NHP obtains the lowest NLL on Taxi and Amazon (on

Amazon by a large margin, $\approx 0,64$ against $\approx 2,2$) and is tied with the THP on Stackoverflow. This result is expected and does not contradict our contribution, for two reasons. First, it agrees with the unified benchmark of the area: when all the models are reimplemented with the same likelihood routine, the recurrent NHP is the strongest on several datasets, winning exactly on Taxi and Retweet in the EasyTPP benchmark (XUE *et al.*, 2024), because its continuous-time formulation models the density between events natively. Second, the NLL is only one of the metrics: on type accuracy the HoTHP surpasses the NHP on Amazon (0,340 against 0,315) and on Stackoverflow (0,451 against 0,445), so the advantage of the NHP is restricted to the likelihood. The research question of this work is not whether attention beats recurrence in absolute NLL, but whether changing the positional kernel improves the extrapolation of a Transformer. For that question the relevant comparison is the controlled one above, in which the HoTHP is consistently the best of the three.

Retweet. This dataset is treated separately in Section 5.4, because with raw timestamps the NLL of the three Transformer models overflows numerically (the extreme temporal scale, with intervals of up to 582.928 time units, makes the hyperbolic and trigonometric kernels saturate, as discussed in Section 4.5). The ablation in Section 5.4 normalizes the scale and isolates the effect of the kernel.

⁹⁸ Figure 12 shows the NLL as a function of the extrapolation factor. We omit the NHP from the plot: its much lower NLL on Taxi and Amazon compresses the vertical scale and hides the differences among the Transformer models, which are the focus of the controlled comparison. The NHP values stay in Table 6. The figure makes the within-family behavior clear: the HoTHP (red) keeps the lowest NLL among the three and the most stable curve, while the RoTHP (blue) shows the widest error band on Taxi at $5\times$ and the only clear divergence on the normalized Retweet.

Figure 13 shows the type-prediction accuracy under the same protocol. On Amazon the recency bias of the HoTHP produces a clear and growing advantage: its accuracy rises to 0,340 at $5\times$, well above the THP and the RoTHP, which stay around 0,31–0,32. On Stackoverflow the HoTHP stays slightly ahead at every factor, and on Taxi the three models are close. As in the NLL plot, the NHP is omitted from the figure to keep the comparison clean within the Transformer family; its accuracy is in Table 7, where the HoTHP also surpasses the NHP on Amazon and Stackoverflow.

Figura 12 – NLL (mean $\pm$ standard deviation, 3 seeds) as a function of the extrapolation factor on the real datasets, for the three Transformer models. The NHP is omitted from the plot for readability, since its lower NLL on Taxi and Amazon would compress the scale; its values are in Table 6. Among the Transformer models, the HoTHP keeps the lowest and most stable NLL in extrapolation.





Figura 13 – Type-prediction accuracy (mean $\pm$ standard deviation, 3 seeds) as a function of the extrapolation factor on the real datasets, for the three Transformer models (the NHP is in Table 7). On Amazon, the HoTHP shows the clearest and growing advantage as the sequence gets longer.



5.2.1 Training convergence


Figure 14 shows the convergence curves of the validation NLL over the 100 training epochs. On the Amazon and Stackoverflow datasets, the three models converge smoothly, with the NLL stabilizing after epoch 50. On Amazon, the curve still shows a slight downward trend in the last 20 epochs, but with a slope below 0,01 per epoch, which indicates that additional epochs

would benefit all the models equally without changing the ranking.

On Taxi, the RoTHP shows higher variance between seeds in the early epochs (10 to 40), but converges to the same level as the HoTHP after epoch 40. On Retweet, the convergence reveals the numerical collapse: 2 out of 3 seeds stabilize at $NLL \approx 15,9$, while only one of the three seeds converges to informative values ($NLL \approx 5,5$). The resulting standard deviation band is enormous, visually confirming the instability reported in Table 8.

Figura 14 – Convergence curves of the validation NLL (mean $\pm$ standard deviation, 3 seeds). The first epochs were omitted to improve the visualization. On Amazon and Stackoverflow, all the models converge smoothly. On Taxi, the RoTHP is more unstable in the early epochs. On Retweet, the wide band of the RoTHP reflects the collapse of 2 of the 3 seeds.

Convergencia do Treino — NLL de Validacao (media $\pm$ std, 3 seeds)



5.3 Intensities learned by the models

To understand why the HoTHP behaves differently on each dataset, we extracted the conditional intensity function $\lambda^*(t)$ that each trained model computes for representative test sequences. Figure 15 shows the total intensity $\lambda^*(t) = \sum_k \lambda_k^*(t)$ over time for the Amazon, Stackoverflow, and Retweet datasets.

On Amazon, the intensity grows between events and falls when an event occurs, indicating that each event reduces the chance of new events in the short term. Even so, the most

Figura 15 – Conditional intensity $\lambda^*(t)$ learned by the THP (orange), RoTHP (blue), and HoTHP (green) for one test sequence of each dataset. The gray vertical bars indicate the events. On Amazon, the intensity grows between events and falls when an event occurs (self-inhibition). On Stackoverflow, the intensity is approximately constant (memoryless process). On Retweet, both models collapse to $\lambda^* \approx 0$ after the initial burst due to the extreme temporal scale.



recent event is the most informative for the prediction, which explains why the recency bias of the HoTHP helps on this dataset. On Stackoverflow, the three models learn a practically flat intensity, with no visible reaction to the events, characteristic of a process with little temporal memory, and no positional bias manages to stand out. On Retweet, the models only react to the initial burst and then the intensity falls to practically zero, because the intervals between events are too large for any decay function.

The observation on Amazon is particularly relevant in light of Section 4.8: the HoTHP wins on this dataset even though the process is self-inhibiting, not self-exciting. This confirms that the recency bias acts on the attention (“which past events are more relevant”) and not on the intensity (“the intensity rises or falls after an event”). On Amazon, the most recent event is the most informative for predicting the next one and the exponential decay in the

attention captures this correctly.

5.4 Ablation: temporal normalization of Retweet

The numerical collapse observed on Retweet raises a question: is the problem the temporal scale or the HoTHP bias itself? To isolate these factors, we ran an ablation in which we normalize the Retweet timestamps before training.

The normalization consists of dividing all the instants of each sequence by the mean of the intervals of that sequence.

After this operation, the mean interval between events becomes approximately 1 in all the sequences, with the maximum interval reduced from 582.928 to approximately 227. The temporal structure (order of the events, clustering patterns, types) is preserved. Only the numerical scale changes.

We trained the RoTHP and the HoTHP on the normalized Retweet with the same hyperparameters (3 seeds, 100 epochs, $L_{\text{train}} = 50$, $L_{5\times} = 250$) and compared with the original results.

Tabela 8 – Ablation: effect of the temporal normalization on Retweet. “Original (raw)” uses the raw timestamps; “Normalized” uses the timestamps divided by the mean of the gaps. Mean $\pm$ standard deviation (3 seeds). The best value in each column of the normalized block is in bold.


Version	Model	NLL $1\times$	NLL $2\times$	NLL $5\times$	Acc $5\times$
Original (raw)	THP	15,94	$3,6 \times 10^3$	$2,3 \times 10^4$	0,491
	RoTHP	15,94	$6,7 \times 10^5$	$7,0 \times 10^5$	0,491
	HoTHP	15,94	$1,2 \times 10^5$	$1,4 \times 10^5$	0,491
Normalized	THP	$-2,180 \pm 0,008$	$-1,227 \pm 0,053$	$0,039 \pm 0,332$	$0,597 \pm 0,002$
	RoTHP	$-2,133 \pm 0,021$	$0,749 \pm 0,618$	$47,563 \pm 8,703$	$0,589 \pm 0,002$
	HoTHP	$-2,139 \pm 0,014$	$-1,275 \pm 0,013$	$-0,651 \pm 0,013$	$0,600 \pm 0,002$

The normalization resolves the numerical collapse unambiguously. With raw timestamps, the NLL collapses to $\approx 15,9$ already at $1\times$ (the same value for the three models, a clear sign of saturation) and grows to the order of 10^4 – 10^5 at $5\times$. With normalized timestamps, the NLL at $1\times$ drops to negative values and the training converges across all the seeds.

The difference between the models appears in the extrapolation, as illustrated in Figure 16. The normalized HoTHP obtains $\text{NLL } 5\times = -0,651 \pm 0,013$, practically invariant between seeds, and it is the only model that keeps a negative NLL in extrapolation. The

THP degrades to $\approx 0,04$ (with high variance, $\pm 0,33$) and the RoTHP degrades much more, to $47,6 \pm 8,7$. The type accuracy follows the same pattern: the normalized HoTHP keeps the best $\text{Acc } 5\times = 0,600 \pm 0,002$, against 0,597 of the THP and 0,589 of the RoTHP.

Figura 16 – NLL on the normalized Retweet (mean $\pm$ standard deviation, 3 seeds). On the left, the $1\times$ condition (training length): both models obtain a similar NLL. On the right, the $5\times$ condition (extrapolation): the HoTHP keeps a stable NLL, while the RoTHP shows high variance due to one seed that diverges.



These results make it possible to separate the two factors that hurt the performance on Retweet. The primary problem was the numerical scale. The normalization benefits all the models drastically. But even with the scale corrected, the HoTHP surpasses both the THP and the RoTHP in extrapolation, confirming that the recency bias is advantageous on this dataset when the temporal scale allows its calibration. The residual degradation of the RoTHP ($\text{NLL } 5\times = 47,6$, far from the HoTHP) reinforces the fragility of the oscillatory kernel in extrapolation, also observed in the synthetic data.

5.5 Discussion

The results on real data show that the recency bias of the HoTHP is beneficial in some scenarios and neutral or harmful in others. The analysis of the learned intensities suggests that the determining factor is not whether the process is self-exciting or self-inhibiting, but whether recent events are more informative than distant ones for the construction of the hidden state.

On Amazon, where the HoTHP obtains the best result, the process is self-inhibiting: the intensity grows with the time since the last event. Even so, the recency bias helps, because

the immediately previous event is the most relevant for predicting the next one. The exponential decay in the attention correctly captures this relationship of decreasing relevance.

On Stackoverflow, where the three models tie, the intensity is practically constant. This indicates a process with little or no temporal memory, in which all the past events are equally (little) informative. The recency bias of the HoTHP does not significantly hurt (the constant intensity shows that the model manages to compensate in the intensity layer), but it also does not help, because there is no temporal structure to exploit.

On Retweet, the problem is twofold: the extreme temporal scale causes numerical collapse, and the structure of viral cascades may not align with the exponential decay. The ablation in Section 5.4 separated these two factors and showed that the primary problem was the scale: with normalization, the HoTHP obtains the best result of the whole benchmark ($NLL\ 5\times = -0,651$), with almost zero variance between seeds. The RoTHP, even with normalization, still degrades sharply in the extrapolation ($NLL\ 5\times = 47,6$).

The HoTHP proves more effective when two conditions are satisfied at the same time: (i) the process has short-range temporal memory, so that recent events are more informative than distant ones; and (ii) the scale of the intervals between events allows the parameter θ' to be calibrated properly during training, without numerical saturation. The Retweet ablation confirms that condition (ii) can be solved by preprocessing, and when it is, the recency bias of the HoTHP proves consistently advantageous.

5.5.1 Computational cost versus benefit in extrapolation

The hyperbolic kernel of the HoTHP involves more operations than the trigonometric kernel of the RoTHP, which is reflected in the training time. Table 9 presents the average training times (100 epochs, 3 seeds) for each combination of dataset and model.

Tabela 9 – Average training time in minutes (mean over 3 seeds, except NHP on Amazon with 2 seeds; NHP was not run on the normalized Retweet). All the models use the same architecture size and training budget (100 epochs).

Dataset	THP	RoTHP	HoTHP	NHP
Taxi	1,9	1,9	2,0	2,7
Stackoverflow	2,5	2,5	3,2	4,9
Amazon	9,4	9,5	12,2	17,7
Retweet (norm)	34,1	34,4	83,2	—

Compared with the THP and the RoTHP, the HoTHP is between $1,07\times$ (Taxi) and

$2,44\times$ (normalized Retweet) slower, the overhead being more pronounced on datasets with long sequences, where the computation of the hyperbolic kernel dominates the time per batch. For datasets of moderate length (Taxi, Stackoverflow, Amazon), the additional cost stays around 30% or below. It is worth noting that the NHP, the strongest baseline in NLL on these datasets, is the most expensive model of all: it is $1,3\times$ to $1,5\times$ slower than the HoTHP (for example, 17,7 against 12,2 minutes on Amazon). The recency-bias kernel therefore adds a moderate cost over the other Transformers while remaining cheaper than the recurrent alternative.

This cost, however, is paid only once during training. At inference time, the trained model is applied to sequences that are potentially much longer than the ones seen in training, exactly the extrapolation scenario. In this context, what matters is not the training speed, but whether the model produces valid results. ⁴⁵ The results in Table 8 illustrate the point clearly: on the normalized Retweet with $5\times$ extrapolation, the RoTHP obtains $\text{NLL} = 47,6 \pm 8,7$, while the HoTHP obtains $\text{NLL} = -0,651 \pm 0,013$. Training the RoTHP $2,4\times$ faster only to then obtain unusable results in extrapolation does not constitute a practical advantage.

REFERÊNCIAS

- DAI, C.; SHAN, H.; SONG, M.; LIANG, D. **HoPE: Hyperbolic Rotary Positional Encoding for Stable Long-Range Dependency Modeling in Large Language Models**. arXiv, 2025. ArXiv:2509.05218 [cs]. Disponível em: <<http://arxiv.org/abs/2509.05218>>.
- DAI, S.; GAO, A.; LI, Z.; DU, Y. **Rotary Position Embedding-Based Transformer Hawkes Process for Event-Type Big Data**. 2026. Disponível em: <<https://www.sciopen.com/article/10.26599/BDMA.2025.9020029>>.
- DALEY, D. J.; VERE-JONES, D. **An Introduction to the Theory of Point Processes: Volume I: Elementary Theory and Methods**. [S.l.]: Springer Science & Business Media, 2003. Google-Books-ID: Ev2iXQKItpUC. ISBN 978-0-387-95541-4.
- DU, N.; DAI, H.; TRIVEDI, R.; UPADHYAY, U.; GOMEZ-RODRIGUEZ, M.; SONG, L. Recurrent Marked Temporal Point Processes: Embedding Event History to Vector. In: **Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining**. New York, NY, USA: Association for Computing Machinery, 2016. (KDD '16), p. 1555–1564. ISBN 978-1-4503-4232-2. Disponível em: <<https://doi.org/10.1145/2939672.2939875>>.
- HAWKES, A. G. Spectra of some self-exciting and mutually exciting point processes. **Biometrika**, v. 58, n. 1, p. 83–90, abr. 1971. ISSN 0006-3444. Disponível em: <<https://doi.org/10.1093/biomet/58.1.83>>.
- MEI, H.; EISNER, J. **The Neural Hawkes Process: A Neurally Self-Modulating Multivariate Point Process**. arXiv, 2017. ArXiv:1612.09328 [cs]. Disponível em: <<http://arxiv.org/abs/1612.09328>>.
- PRESS, O.; SMITH, N.; LEWIS, M. Train Short, Test Long: Attention with Linear Biases Enables Input Length Extrapolation. In: . [s.n.], 2021. Disponível em: <<https://openreview.net/forum?id=R8sQPpGCv0>>.
- RASMUSSEN, J. G. **Lecture Notes: Temporal Point Processes and the Conditional Intensity Function**. arXiv, 2018. ArXiv:1806.00221 [stat]. Disponível em: <<http://arxiv.org/abs/1806.00221>>.
- ROSS, S. M. **A first course in probability**. Tenth edition. Boston: Pearson, 2019. ISBN 978-0-13-475311-9.
- SHCHUR, O.; TÜRKMEN, A. C.; JANUSCHOWSKI, T.; GÜNNEMANN, S. **Neural Temporal Point Processes: A Review**. arXiv, 2021. ArXiv:2104.03528 [cs]. Disponível em: <<http://arxiv.org/abs/2104.03528>>.
- SU, J.; LU, Y.; PAN, S.; MURTADHA, A.; WEN, B.; LIU, Y. **RoFormer: Enhanced Transformer with Rotary Position Embedding**. arXiv, 2023. ArXiv:2104.09864 [cs]. Disponível em: <<http://arxiv.org/abs/2104.09864>>.
- VASWANI, A.; SHAZEER, N.; PARMAR, N.; USZKOREIT, J.; JONES, L.; GOMEZ, A. N.; KAISER, u.; POLOSUKHIN, I. Attention is All you Need. In: GUYON, I.; LUXBURG, U. V.; BENGIO, S.; WALLACH, H.; FERGUS, R.; VISHWANATHAN, S.; GARNETT, R. (Ed.). **Advances in Neural Information Processing Systems**. Curran Associates, Inc.,

2017. v. 30. Disponível em: <https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf>.

⁴ XUE, S.; SHI, X.; CHU, Z.; WANG, Y.; HAO, H.; ZHOU, F.; JIANG, C.; PAN, C.; ZHANG, J. Y.; WEN, Q.; ⁸ ZHOU, J.; MEI, H. **EasyTPP: Towards Open Benchmarking Temporal Point Processes**. arXiv, 2024. ArXiv:2307.08097 [cs]. Disponível em: <<http://arxiv.org/abs/2307.08097>>.

¹ ZHANG, Q.; LIPANI, A.; KIRNAP, O.; YILMAZ, E. Self-Attentive Hawkes Process. In: **Proceedings of the 37th International Conference on Machine Learning**. PMLR, 2020. p. 11183–11193. ISSN 2640-3498. Disponível em: <<https://proceedings.mlr.press/v119/zhang20q.html>>.

¹ ZHOU, F.; KONG, Q.; QIAO, J.; WAN, C.; ZHANG, Y.; CAI, R. **Advances in Temporal Point Processes: Bayesian, Neural, and LLM Approaches**. arXiv, 2025. ArXiv:2501.14291 [cs]. Disponível em: <<http://arxiv.org/abs/2501.14291>>.

²⁰ ZUO, S.; JIANG, H.; LI, Z.; ZHAO, T.; ZHA, H. **Transformer Hawkes Process**. ⁸ arXiv, 2021. ArXiv:2002.09291 [cs]. Disponível em: <<http://arxiv.org/abs/2002.09291>>.

● 16% geral de similaridade

As principais fontes encontradas nos seguintes bancos de dados:

- 13% Banco de dados da Internet
- Banco de dados do Crossref
- 11% Banco de dados de trabalhos enviados
- 11% Banco de dados de publicações
- Banco de dados de conteúdo publicado no Crossref

PRINCIPAIS FONTES

As fontes com o maior número de correspondências no envio. Fontes sobrepostas não serão exibidas.

1	arxiv.org Internet	2%
2	jhir.library.jhu.edu Internet	1%
3	export.arxiv.org Internet	1%
4	raw.githubusercontent.com Internet	<1%
5	dx.doi.org Internet	<1%
6	proceedings.mlr.press Internet	<1%
7	repositorio.ufc.br Internet	<1%
8	Marcel Rodrigues de Barros. "Incorporando as irregularidades de dado..." Crossref posted content	<1%

9	Alana da Silva Luna. "AllianceScan: uma abordagem para identificação..."	<1%
	Crossref posted content	
10	mdpi.com	<1%
	Internet	
11	University of Utah on 2026-04-24	<1%
	Submitted works	
12	arxiv-vanity.com	<1%
	Internet	
13	jku on 2025-07-23	<1%
	Submitted works	
14	escholarship.org	<1%
	Internet	
15	Özyegin Üniversitesi on 2026-06-16	<1%
	Submitted works	
16	Universidade Federal do Ceará, UFC on 2022-11-01	<1%
	Submitted works	
17	University College Roosevelt on 2025-07-28	<1%
	Submitted works	
18	Aston University on 2023-09-28	<1%
	Submitted works	
19	ebin.pub	<1%
	Internet	
20	frontiersin.org	<1%
	Internet	

21	Yin Li. "Theoretical Analysis of Positional Encodings in Transformer M... Crossref posted content	<1%
22	discovery-pp.ucl.ac.uk Internet	<1%
23	sh-tsang.medium.com Internet	<1%
24	kar.kent.ac.uk Internet	<1%
25	"Machine Learning and Knowledge Discovery in Databases: Research T... Crossref	<1%
26	fiveable.me Internet	<1%
27	"ECAI 2020", IOS Press, 2020 Crossref	<1%
28	Ali Khodadadi, Seyed Abbas Hosseini, Erfan Tavakoli, Hamid R. Rabiee.... Crossref	<1%
29	Indian Institute of Technology, Bombay on 2019-04-08 Submitted works	<1%
30	pleasedameunkiss.blogspot.com Internet	<1%
31	research-collection.ethz.ch Internet	<1%
32	Symbiosis International University on 2017-09-20 Submitted works	<1%

33	houseofwisdom.online	Internet	<1%
34	scielo.br	Internet	<1%
35	Angelica Liguori, Luciano Caroprese, Marco Minici, Bruno Veloso, Fran...	Crossref	<1%
36	Imperial College of Science, Technology and Medicine on 2021-09-03	Submitted works	<1%
37	Massey University on 2017-03-14	Submitted works	<1%
38	University of Warwick on 2016-05-17	Submitted works	<1%
39	expeditiorepositorio.utadeo.edu.co	Internet	<1%
40	math.mit.edu	Internet	<1%
41	Indian Institute of Science, Bangalore on 2024-06-21	Submitted works	<1%
42	Universiteit Hasselt on 2026-06-16	Submitted works	<1%
43	Utkarsh, Kumar. "Reliability of Model Selection in Temporal Data.", Nort...	Publication	<1%
44	ynu.repo.nii.ac.jp	Internet	<1%

45	"Advanced Intelligent Computing Technology and Applications", Spring...	<1%
	Crossref	
46	Indian Institute of Technology on 2022-06-29	<1%
	Submitted works	
47	Qi Song, Wenjie Luo. "SFBKT: A Synthetically Forgetting Behavior Meth...	<1%
	Crossref	
48	University College London on 2018-09-09	<1%
	Submitted works	
49	Ain Shams University on 2020-08-23	<1%
	Submitted works	
50	University College London on 2018-09-03	<1%
	Submitted works	
51	playbigdata.ruc.edu.cn	<1%
	Internet	
52	politesi.polimi.it	<1%
	Internet	
53	Birla Institute of Technology and Science Pilani on 2021-06-21	<1%
	Submitted works	
54	University of Newcastle upon Tyne on 2013-12-10	<1%
	Submitted works	
55	Imperial College of Science, Technology and Medicine on 2022-09-02	<1%
	Submitted works	
56	Lei Jiang, Huan Liu, Jesse S. Jin, Yu Zheng, Xin He, Jie Zhao. "TransSe...	<1%
	Crossref	

57	Pereira, Kleiton. "Scheduling HPC Jobs with Graph Neural Networks", U... Publication	<1%
58	United International University on 2021-03-12 Submitted works	<1%
59	University of Leeds on 2025-08-26 Submitted works	<1%
60	artemis.cslab.ece.ntua.gr:8080 Internet	<1%
61	proceedings.iclr.cc Internet	<1%
62	"Combinatorial Optimization and Applications", Springer Science and B... Crossref	<1%
63	Heriot-Watt University on 2023-11-29 Submitted works	<1%
64	Liverpool John Moores University on 2023-12-04 Submitted works	<1%
65	Siawpeng Er, Shihao Yang, Tuo Zhao. "COUnty aggRegation mixup AuG... Crossref	<1%
66	sol.sbc.org.br Internet	<1%
67	Coventry University on 2023-09-08 Submitted works	<1%
68	Indian Institute of Science, Bangalore on 2025-06-25 Submitted works	<1%

69	Lars Schäfer. "<mml:math altimg='si1.gif' overflow='scroll' xmlns..."	<1%
	Crossref	
70	Tatsunori Tanaka, Fi Zheng, Kai Sato, Zhifeng Li, Shi Li. "Time-Aware P..."	<1%
	Crossref posted content	
71	Universidade do Porto on 2019-07-22	<1%
	Submitted works	
72	University of Birmingham on 2021-09-13	<1%
	Submitted works	
73	University of Sheffield on 2025-06-15	<1%
	Submitted works	
74	alumni-portal.sasin.edu	<1%
	Internet	
75	huggingface.co	<1%
	Internet	
76	mason.gmu.edu	<1%
	Internet	
77	repo.library.stonybrook.edu	<1%
	Internet	
78	Debanjan Banerjee, Anindita Banerji, Arunita Banerji. "Chapter 2 Appeal..."	<1%
	Crossref	
79	Enzo Orsingher, Riccardo Cesari, Vieri Mosco. "Poisson Process and it..."	<1%
	Publication	
80	Feng , Liang. "Importance Sampling of Affine Jump Diffusions and Rela..."	<1%
	Publication	

81	François Delarue, Mario Maurelli. "Zero Noise Limit for Multidimension...	Crossref	<1%
82	Hojjat Karami, David Atienza, Anisoara Ionescu. "TEE4EHR: Transform...	Crossref	<1%
83	Imperial College of Science, Technology and Medicine on 2026-06-12	Submitted works	<1%
84	Indian Institute of Technology on 2022-07-24	Submitted works	<1%
85	P.N. Paraskevopoulos. "Modern Control Engineering", CRC Press, 2017	Publication	<1%
86	Universiteit van Amsterdam on 2018-09-28	Submitted works	<1%
87	University of Edinburgh on 2024-04-25	Submitted works	<1%
88	University of Warwick on 2014-05-01	Submitted works	<1%
89	Y.K. Cheung, J.H.W. Lee, A.Y.T. Leung. "Computational Mechanics, Vol...	Publication	<1%
90	cdr.lib.unc.edu	Internet	<1%
91	hal.science	Internet	<1%
92	jscholarship.library.jhu.edu	Internet	<1%

93	papers.nips.cc Internet	<1%
94	unifei on 2024-08-28 Submitted works	<1%
95	hotpaper.io Internet	<1%
96	"Machine Learning and Knowledge Discovery in Databases. Research T... Crossref	<1%
97	Bocconi University on 2011-10-28 Submitted works	<1%
98	C. Guedes Soares, F.C. Salvado. "Innovations in Maritime Technology a... Publication	<1%
99	Erasmus University Rotterdam on 2026-06-20 Submitted works	<1%
100	Heriot-Watt University on 2024-04-12 Submitted works	<1%
101	Imperial College of Science, Technology and Medicine on 2020-10-12 Submitted works	<1%
102	Imperial College of Science, Technology and Medicine on 2022-06-22 Submitted works	<1%
103	Imperial College of Science, Technology and Medicine on 2026-06-02 Submitted works	<1%
104	Indian Institute of Technology on 2019-06-20 Submitted works	<1%

105	Laila F. Seddek, Abdelhalim Ebaid, Essam R. El-Zahar, Mona D. Aljoufi. ... Crossref	<1%
106	Louise Grinstein, Sally I. Lipsey. "Encyclopedia of Mathematics Educati... Publication	<1%
107	Lu-ning Zhang, Jian-wei Liu, Zhi-yan Song, Xin Zuo. "Temporal attentio... Crossref	<1%
108	Mingxu Tao, Yansong Feng, Dongyan Zhao. "Chapter 24 A Frustratingly... Crossref	<1%
109	Qiang Zhang, Aldo Lipani, Emine Yilmaz. "Learning Neural Point Proces... Crossref	<1%
110	Queen Mary and Westfield College on 2023-08-24 Submitted works	<1%
111	Shengran Guo. "Application and Impact of Position Embedding in Natu... Crossref	<1%
112	Shuting Li, Shuanghong Shen, Yu Su, Xinjie Sun, Junyu Lu, Qi Mo, Zhen... Crossref	<1%
113	University College London on 2025-04-27 Submitted works	<1%
114	University of Birmingham on 2025-08-30 Submitted works	<1%
115	University of Leicester on 2026-05-27 Submitted works	<1%
116	downloads.hindawi.com Internet	<1%

117	ojs.aaai.org Internet	<1%
118	pure.uva.nl Internet	<1%
119	roots-automation.github.io Internet	<1%
120	coursehero.com Internet	<1%
121	lrec-conf.org Internet	<1%
122	"An Introduction to the Theory of Point Processes", Springer Nature, 20... Crossref	<1%
123	Higher Education Commission Pakistan on 2016-11-18 Submitted works	<1%
124	Liang Chen, Shyuichi Izumiya, DongHe Pei, Kentaro Saji. "Anti de Sitter ... Crossref	<1%
125	Oklahoma State University on 2015-11-07 Submitted works	<1%
126	Pierre Del Moral, Spiridon Penev. "Stochastic Processes - From Applic... Publication	<1%
127	University of Birmingham on 2020-09-03 Submitted works	<1%
128	University of Lancaster on 2026-04-24 Submitted works	<1%

-
- 129** Xu, Zhiqing. "A General-Purpose Deep Learning Framework for Protein ... **<1%**
Publication
-
- 130** Monteiro, Andreia Alves Forte Oliveira. "Contributions to Spatial and Te... **<1%**
Publication
-
- 131** Siawpeng Er, Shihao Yang, Tuo Zhao. "County augmented transformer ... **<1%**
Crossref
-
- 132** University of Warwick on 2019-09-13 **<1%**
Submitted works
-
- 133** Université Sétif 1 Ferhat Abbas on 2026-06-14 **<1%**
Submitted works
-
- 134** Zhou, Zihao. "Neural Point Process for Learning Spatiotemporal Event ... **<1%**
Publication